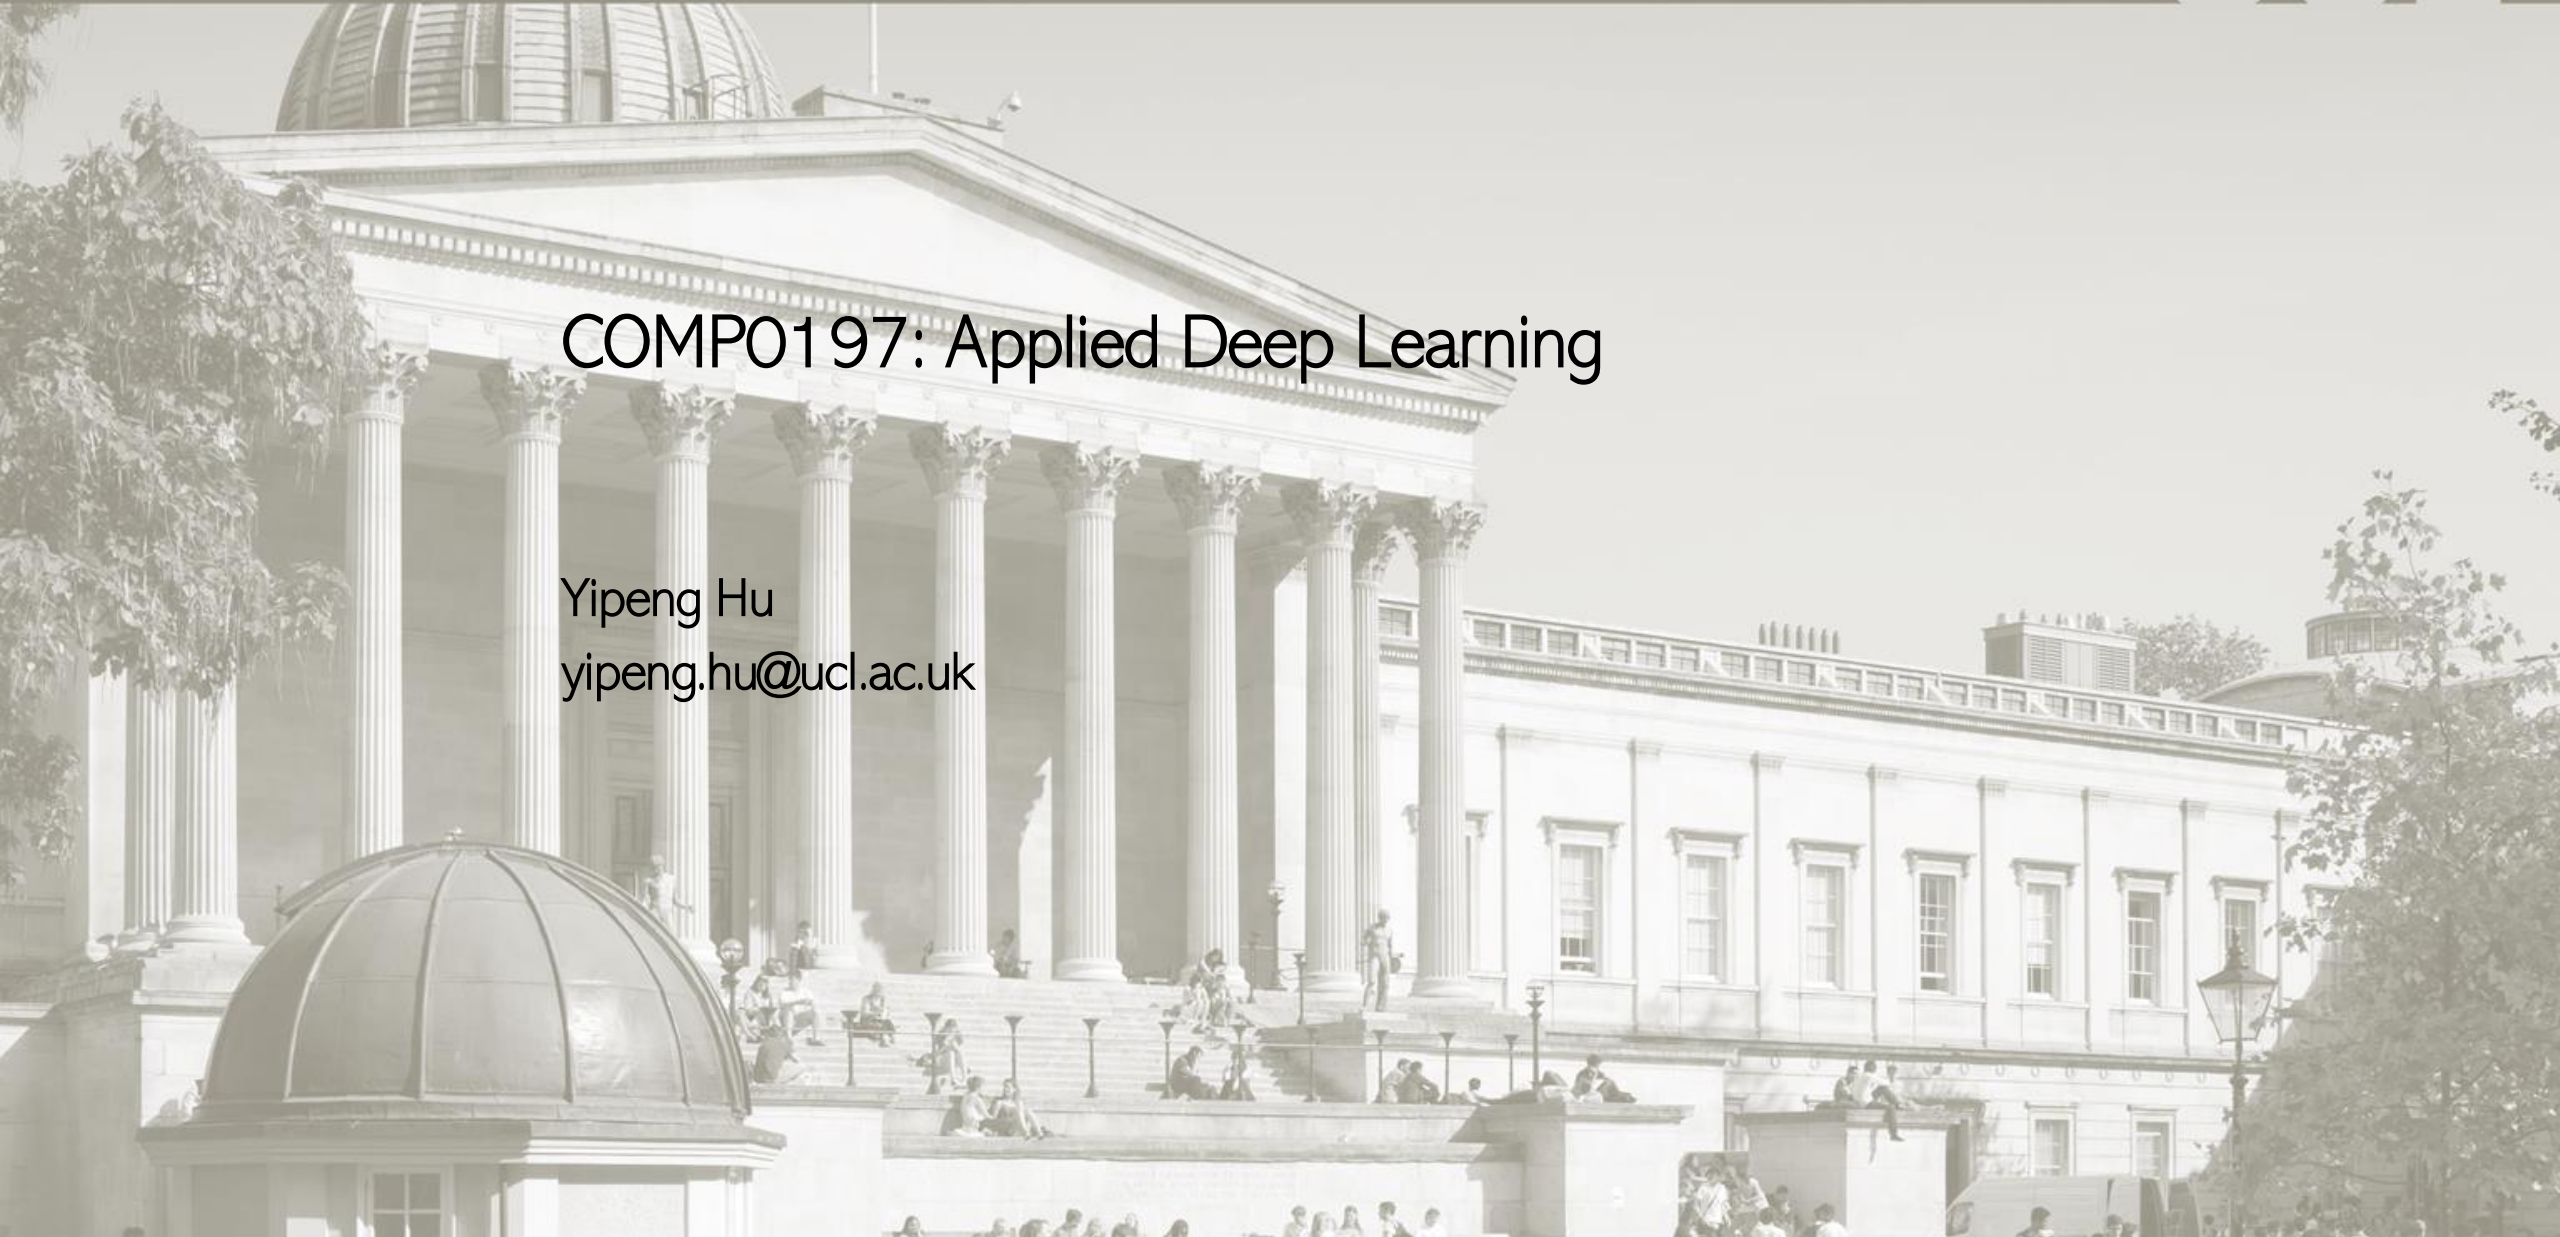


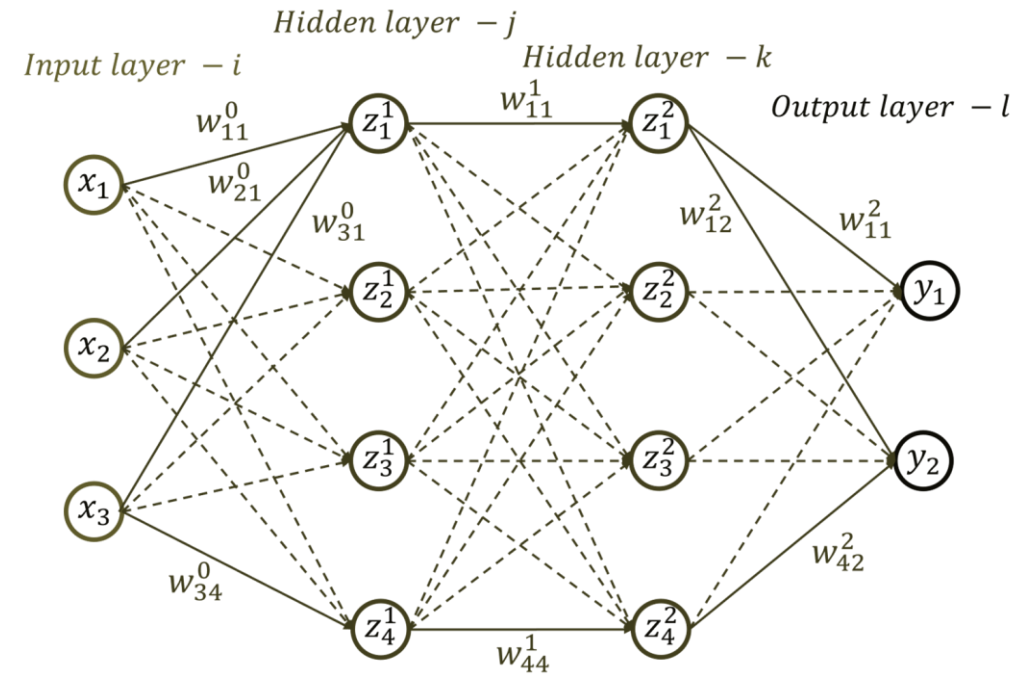
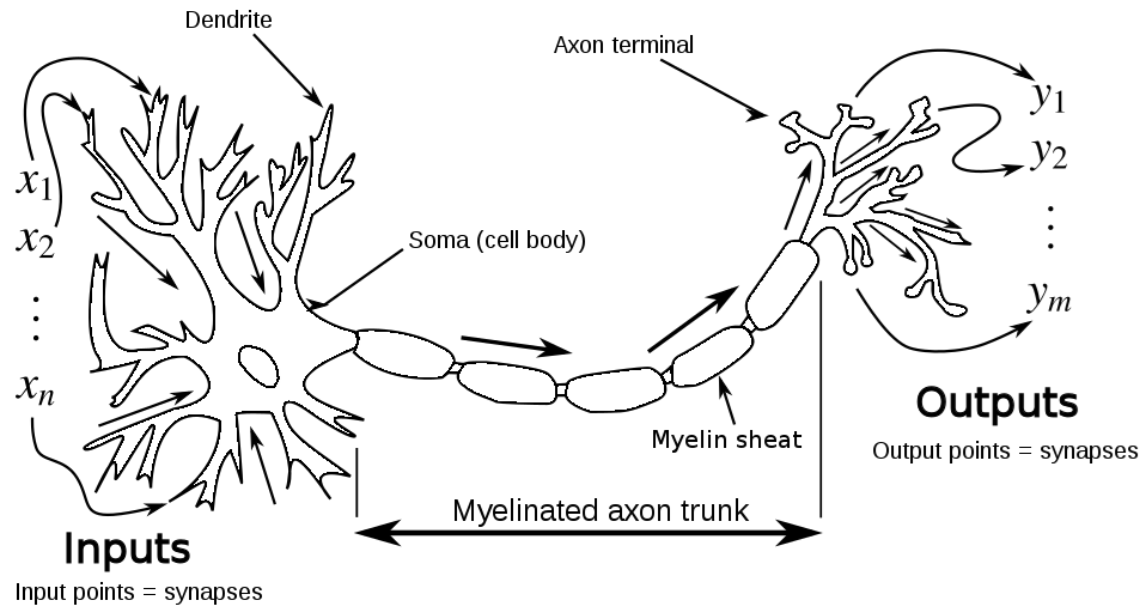
# COMP0197: Applied Deep Learning

Yipeng Hu  
[yipeng.hu@ucl.ac.uk](mailto:yipeng.hu@ucl.ac.uk)





# Deep Feedforward Neural Networks

## Biological and artificial neurons




## Feedforward models



- Multilayer Perceptron (MLP)
- Anatomy of A Layer
- Activation Functions
- Architecture


Width and Depth | Branching and Joining | Skipping | Sampling | “Ignoring”

- Architecture Search



# Deep Feedforward Neural Networks | MLP

## Logistic regression




$x$ : Input

$y$ : Output

$\mathbf{w}$ : Parameters/weights

$$y = \sigma \left[ \sum_{i=1}^{i=3} x_i w_i + b \right] = \sigma[\mathbf{w}^T \mathbf{x} + b]$$

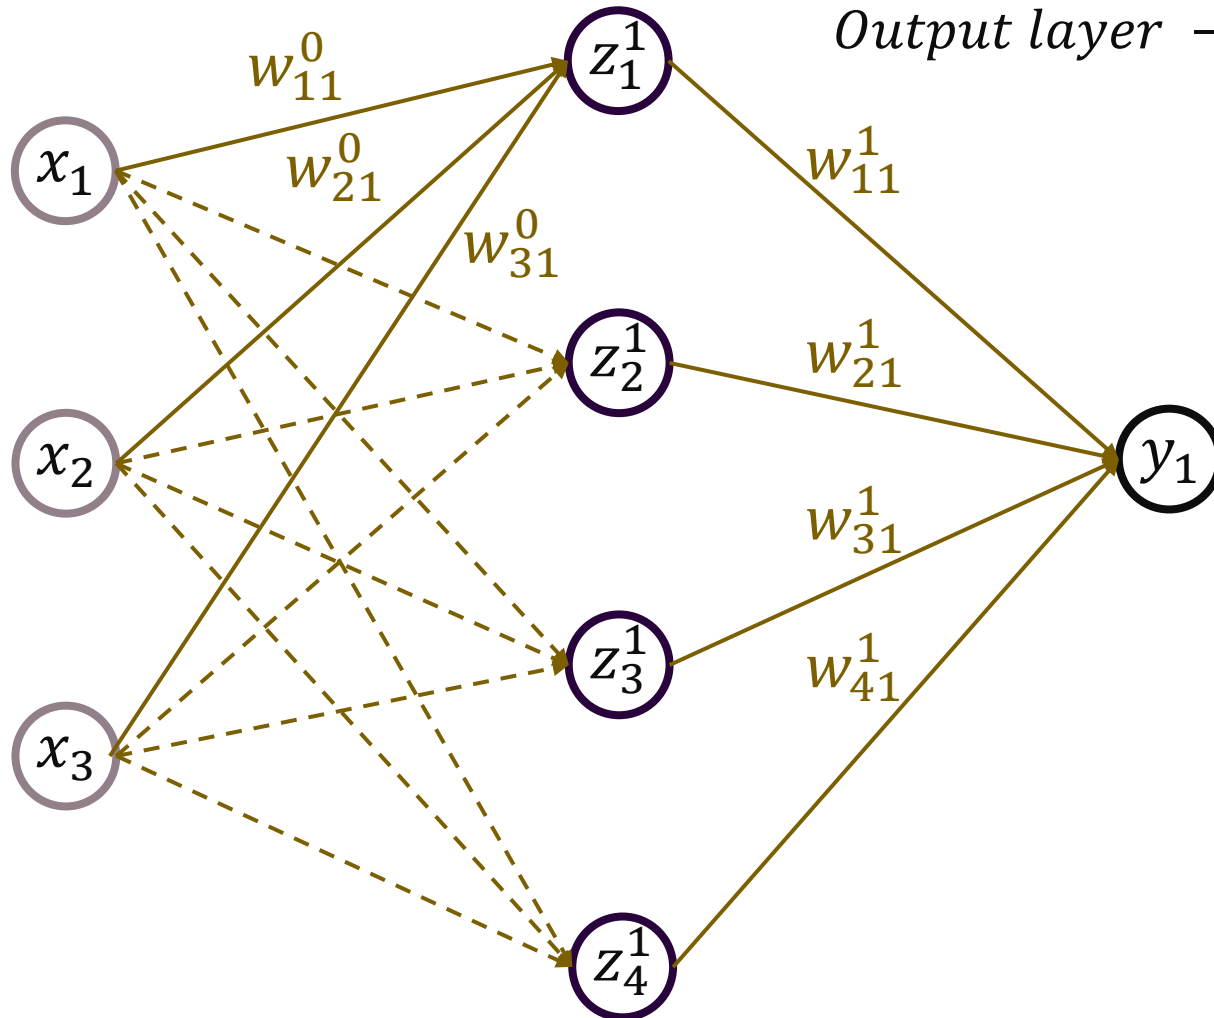
Activation function:  $\sigma(z) = \frac{1}{1 + \exp(-z)}$



Input layer –  $i$

Hidden layer –  $j$

Output layer –  $k$



Activation function:  $\sigma \rightarrow g$

Input  $\rightarrow$  H1:  $i \rightarrow j$


$$z_j^1 = g_j^1 \left[ \sum_{i=1}^{i=3} x_i w_{ij}^0 + b_j^1 \right] = g_j^1 \left[ (\mathbf{w}_j^0)^T \mathbf{x} + b_j^1 \right]$$

H1  $\rightarrow$  Output:  $j \rightarrow k$

$$f(\mathbf{x}, \mathbf{W}) = y_k = g_k \left[ (\mathbf{w}_k^1)^T \mathbf{z}^1 + b_k \right]$$

where  $\mathbf{z}^1 = [z_1^1, z_2^1, z_3^1, z_4^1]^T, k = 1$





Input  $\rightarrow$  H1:  $i \rightarrow j$

$$z_j^1 = g_j^1 \left[ (\mathbf{w}_j^0)^T \mathbf{x} + b_j^1 \right]$$

H1  $\rightarrow$  H2:  $j \rightarrow k$

$$z_k^2 = g_k^2 \left[ (\mathbf{w}_k^1)^T \mathbf{z}^1 + b_k^2 \right]$$

H2  $\rightarrow$  Output:  $k \rightarrow l$

$$y_l = g_l \left[ (\mathbf{w}_l^2)^T \mathbf{z}^2 + b_l^2 \right]$$

# Deep Feedforward Neural Networks | Anatomy of A Layer

$$h = g(\mathbf{w}^T \mathbf{x} + b)$$

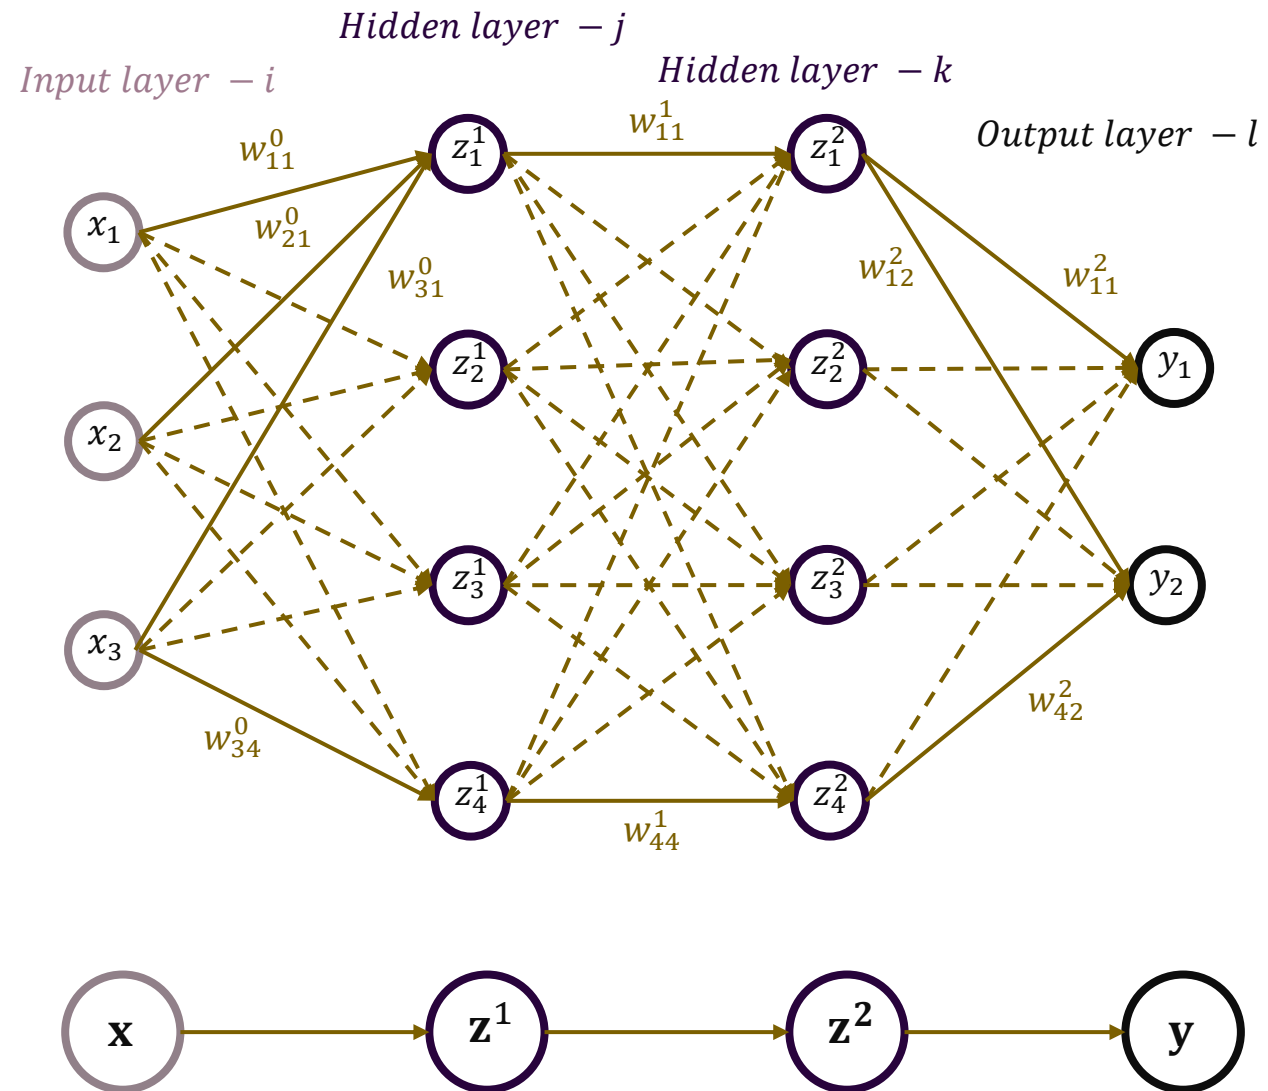
$h$ : Hidden layer & output feature (vectors)

$\mathbf{x}$ : Input feature vector

$\mathbf{w}$ : Weights – network parameters


$b$ : Bias

$g(\cdot)$ : Activation function




# Deep Feedforward Neural Networks | Activation Functions

Nonlinearity?



## Nonlinearity

$$\begin{aligned}y &= \mathbf{w}_j^T (\mathbf{w}_i^T \mathbf{x} + \mathbf{b}_i) + \mathbf{b}_j \\&= \mathbf{w}_j^T \mathbf{w}_i^T \mathbf{x} + (\mathbf{w}_j^T \mathbf{b}_i + \mathbf{b}_j) \\&= \tilde{\mathbf{w}}^T \mathbf{x} + \tilde{\mathbf{b}}\end{aligned}$$




Nonlinearity

Range **[0, 1]**

The squashing effect

A “unit” usually has one activation value



$$g(z) = \frac{1}{1 + e^{-z}}$$

The sigmoid function

## For hidden units

Logistic sigmoid

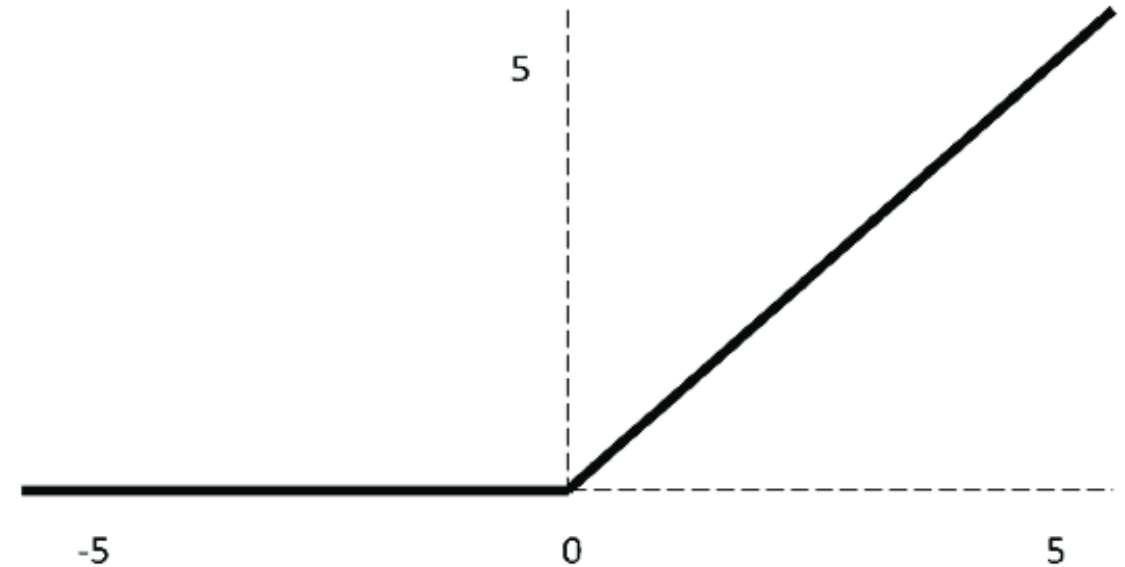

$$g(z) = \sigma(z) = \frac{1}{1+e^{-z}}$$

Hyperbolic tangent

$$g(z) = \tanh(z) = 2\sigma(2z) - 1$$

Rectified linear unit (ReLU)

$$g(z) = \max\{0, z\}$$

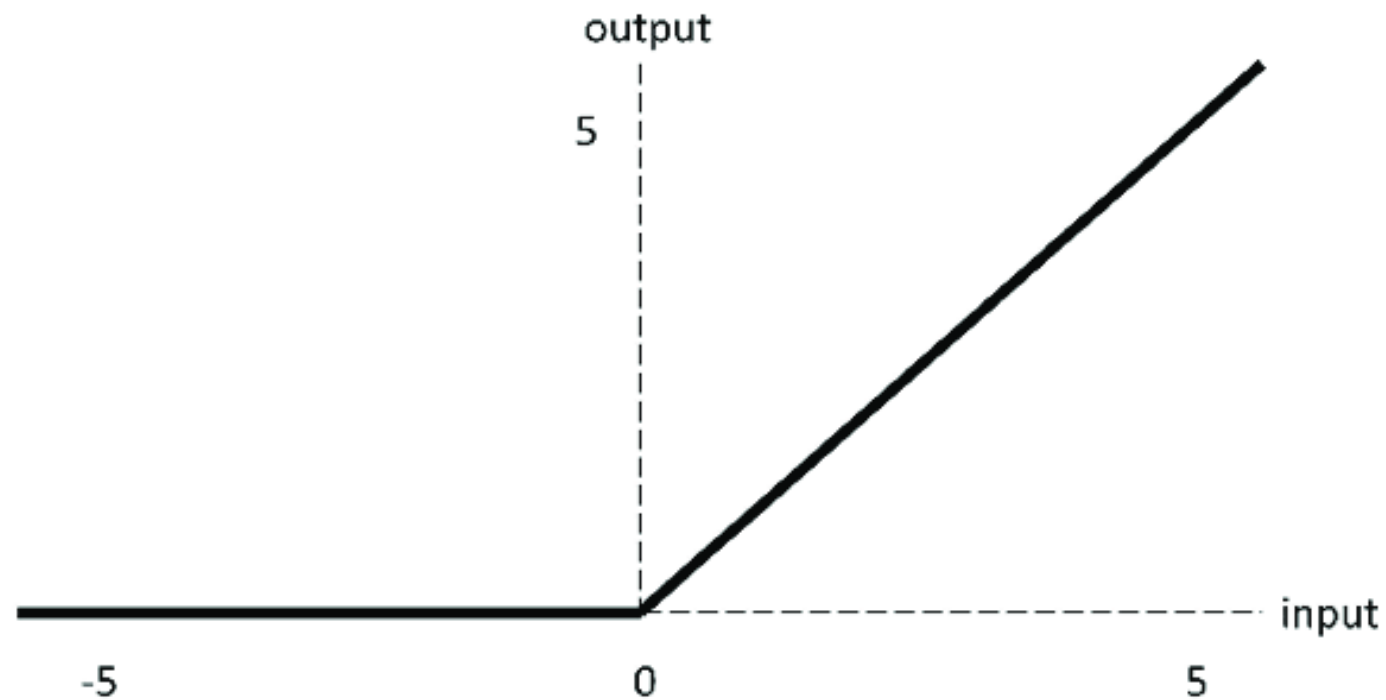




For hidden units

Rectified linear unit (ReLU)

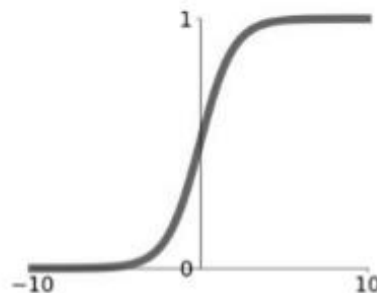
$$g(z) = \max\{0, z\}$$



For hidden units

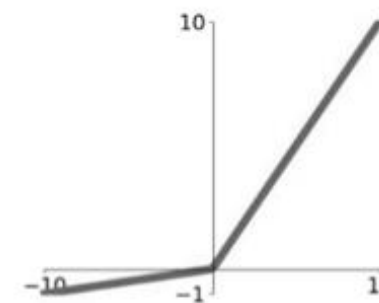
## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



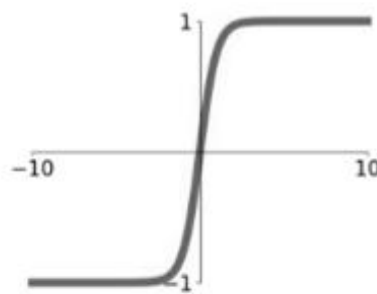
## Leaky ReLU

$$\max(0.1x, x)$$



## tanh

$$\tanh(x)$$

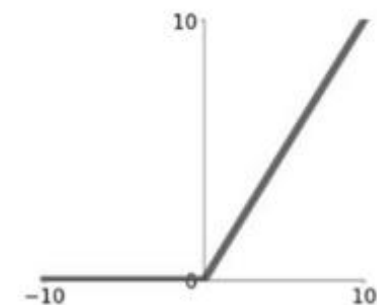


## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

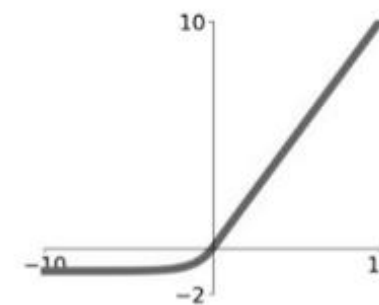
## ReLU

$$\max(0, x)$$



## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$





For output units

Bernoulli output (logistic sigmoid):  $y = g(z) = \frac{1}{1+e^{-z}}$

Multinoulli output (softmax):  $y_c = g(z_c) = \frac{\exp(z_c)}{\sum_{j=1}^C \exp(z_j)}$

$$\rightarrow \begin{cases} \sum_{j=1}^C y_c = 1 \\ 0 \leq y_c \leq 1 \end{cases}$$

Gaussian output



## Other activation functions

Permutation-invariant, e.g. max, mean, min...

Known specific range and distribution

e.g. displacement/velocity (optical flow, registration), function parameters (camera position, pose)

# Deep Feedforward Neural Networks | Architecture

Width and depth


How many layers?

Universal approximation theorem

A sufficiently-wide network approximates any “practically useful” functions

The wider the better?



The deeper the better?



## Width and depth

## Why deep?

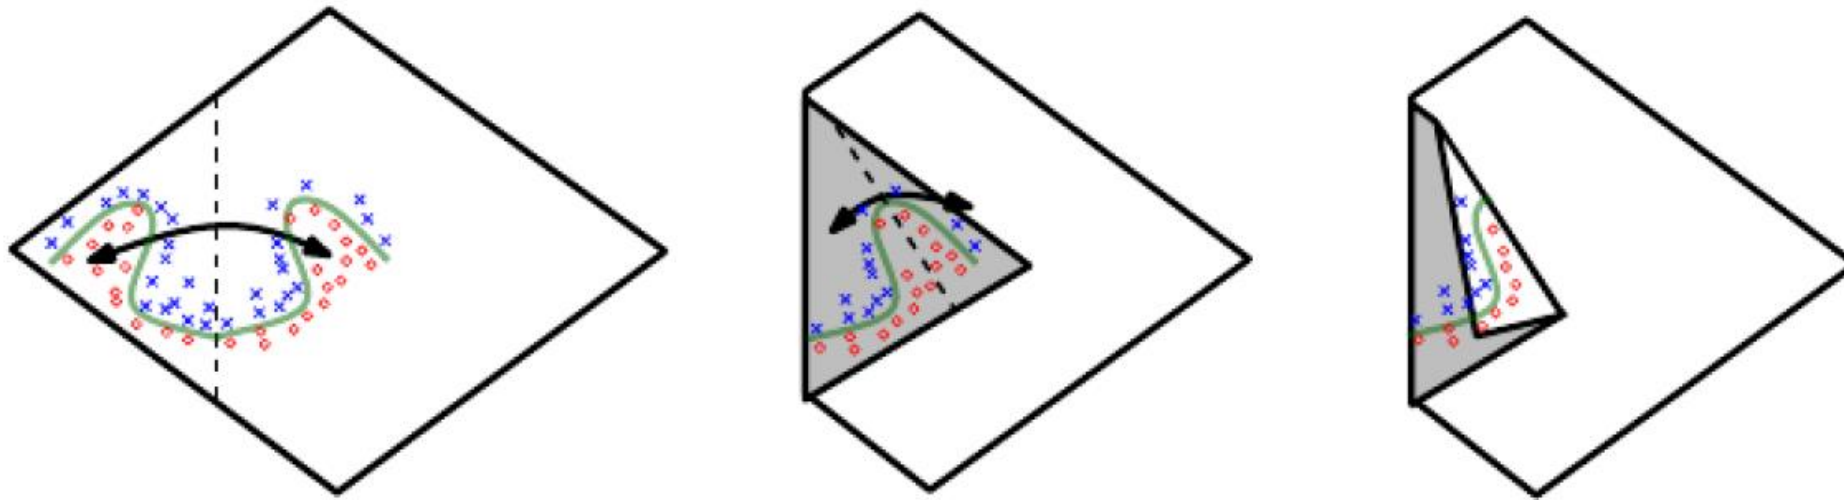
Hierarchical representation learning, Empirical results



## Width and depth

Why deep?

Hierarchical representation learning, Empirical results, Efficiency







## Size

Number of parameters vs storage size ?

Memory consumption ?

Trade-off between accuracy and size

- “MobileNet” & “EfficientNet”\*



## Branching and joining

- Multi-stream

- Multi-task\*

- “Inception”

- Types of input

- Types of output/loss

- Types of processing

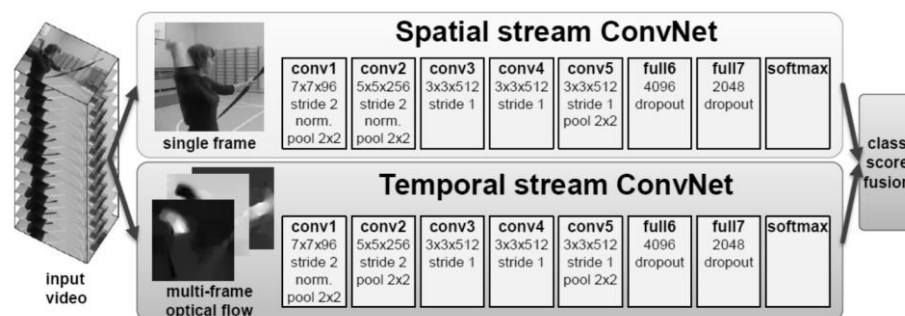




Figure 1: Two-stream architecture for video classification.



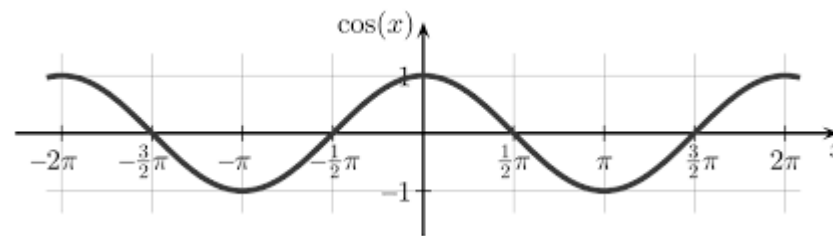
## Branching and joining

- Implementing branching with two sets of (non-)shared weights:

$$h_1 = g(\mathbf{w}_1^T \mathbf{x} + b_1) \text{ and } h_2 = g(\mathbf{w}_2^T \mathbf{x} + b_2)$$

- Implementing joining:

Dot product:  $\mathbf{h}_1 \cdot \mathbf{h}_2 = |\mathbf{h}_1| |\mathbf{h}_2| \cos \theta$  ?




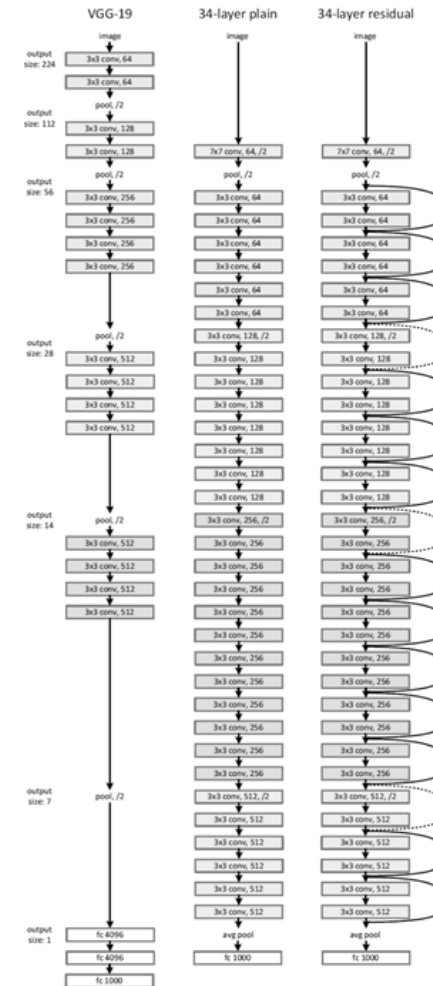
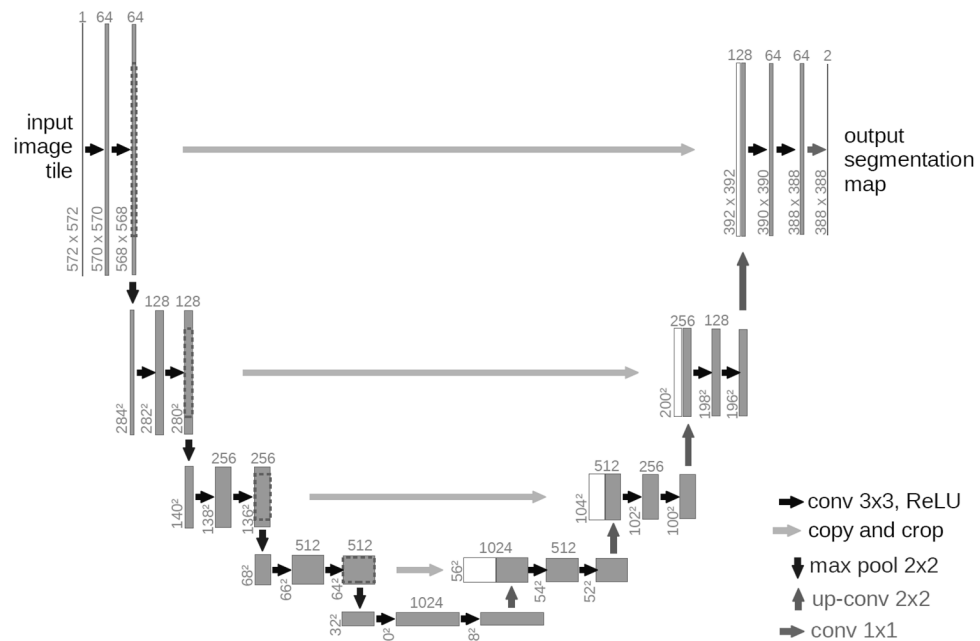
Concatenation vs. summation



$$(\mathbf{h}_1, \mathbf{h}_2) \rightarrow [\mathbf{h}_1^T, \mathbf{h}_2^T]^T$$

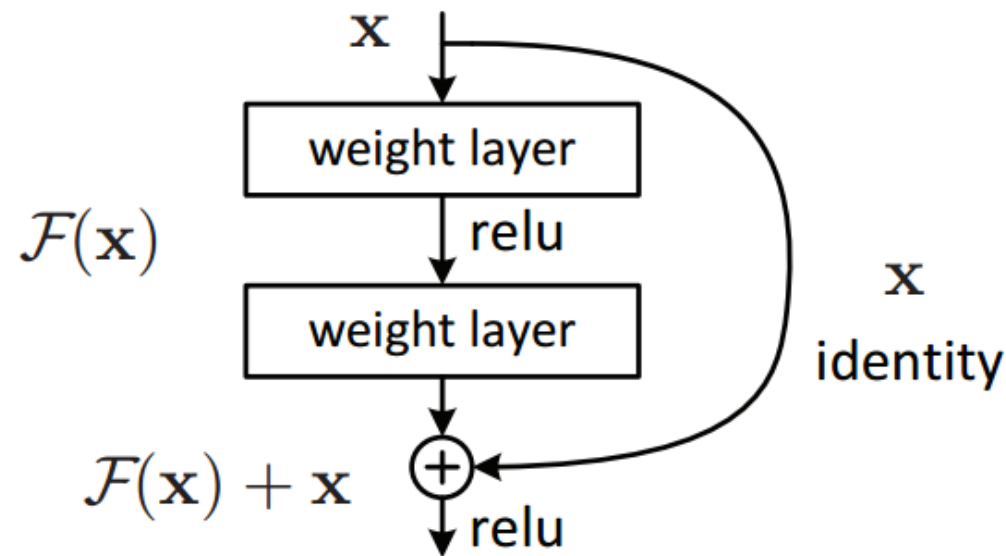
$$(\mathbf{h}_1, \mathbf{h}_2) \rightarrow \mathbf{h}_1 + \mathbf{h}_2$$

- Shortcuts, skip layers, residual connections
- “U-Net”\*



## Skipping

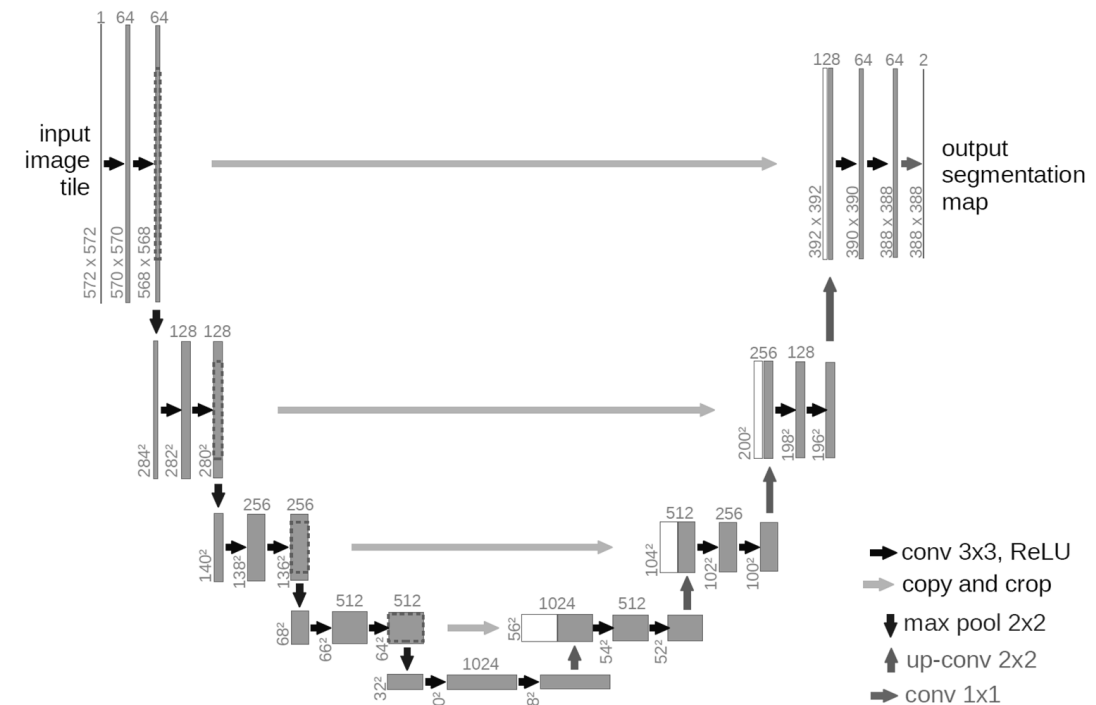


- Shortcuts, skip layers, residual connections
- “U-Net”\*
- “ResNet”\*



$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{W_i\}) + \mathbf{x}.$$

## Sampling

- Down-sampling for encoding ?
- “AlexNet” & “VGG”
- Up-sampling for decoding ?
- “U-Net” & “Autoencoder”
- Sampling for convolution\*



## Sampling

- Implement vector down-sampling: averaging, matrix multiplication
- Implement vector up-sampling: matrix multiplication, interpolation ?

$$h_n = g(\mathbf{w}_n^T \mathbf{x} + b_n)$$

$$\begin{bmatrix} h_1 \\ \vdots \\ h_N \end{bmatrix} = g \left( \begin{bmatrix} \mathbf{w}_1^T \mathbf{x} + b_1 \\ \vdots \\ \mathbf{w}_N^T \mathbf{x} + b_N \end{bmatrix} \right) = g([\mathbf{w}_1 \quad \dots \quad \mathbf{w}_N]^T \mathbf{x} + [b_1 \quad \dots \quad b_N]^T)$$

$$\mathbf{h}^{(N \times 1)} = \mathbf{W}^{(N \times M)} \mathbf{x}^{(M \times 1)} + \mathbf{b}^{(N \times 1)}$$

- Implement feature map sampling\*

“Ignoring”

$$f(\{x_1, x_2, x_3, \dots x_n\}) \approx g(\{h(x_1), h(x_2), h(x_3), \dots h(x_n)\})$$



## “Ignoring”

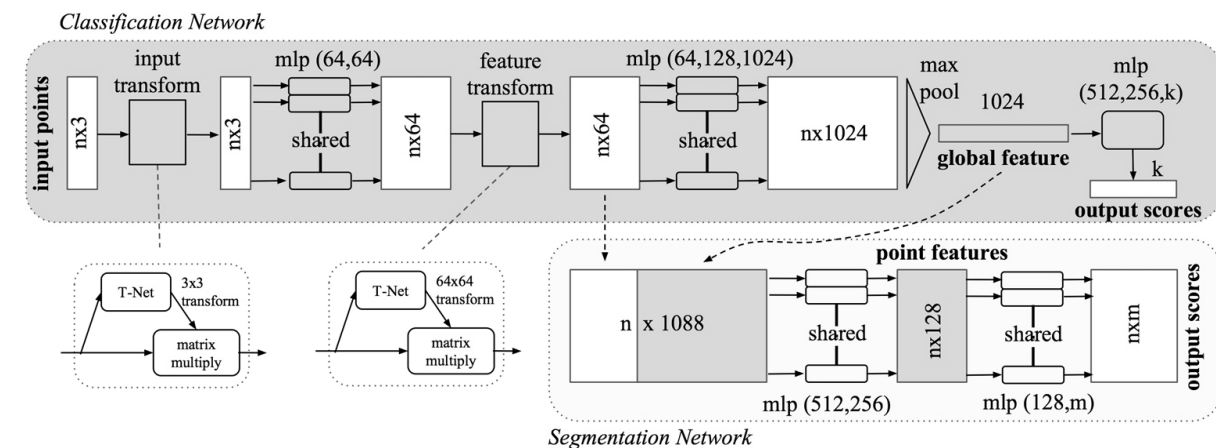


- Permutation-invariant

$$f(\{x_1, x_2, x_3, \dots, x_n\}) \approx g(\{h(x_1), h(x_2), h(x_3), \dots, h(x_n)\})$$

- “PointNet”

max, min, mean...

- Translation-invariant\* (pooling, convolution)



# Deep Feedforward Neural Networks | Architecture Search

## Hyperparameters

- Predefined
- Might not optimisable
- Should not be optimised
- e.g. degree of polynomial, NN architecture choice (width, depth, branches...)

## Hyperparameter tuning

## Hyperparameter search



Training-validation-test 3-way split

Grid search vs. random search

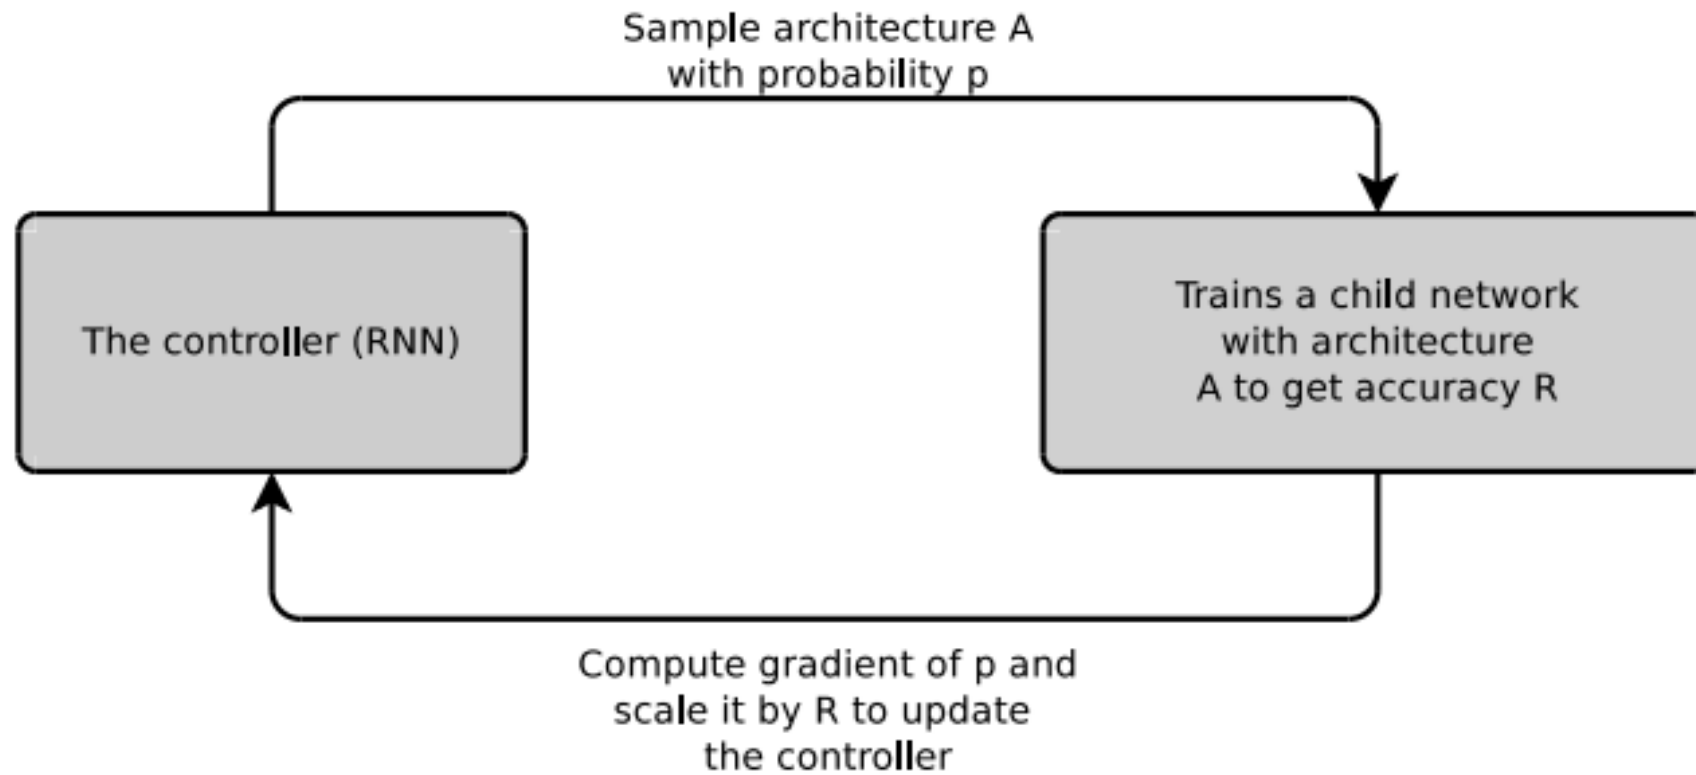
As an optimisation,

e.g. evolution algorithms

- “EfficientNet”\*







## Meta-learning | reinforcement learning



## Meta-learning | differentiable architecture search


- “DARTS”




- Multilayer Perceptron (MLP)
- Anatomy of A Layer
- Activation Functions
- Architecture

Width and Depth | Branching and Joining | Skipping | Sampling | “Ignoring”

- Architecture Search





Using a fully-connected network for the “image classification” tutorial