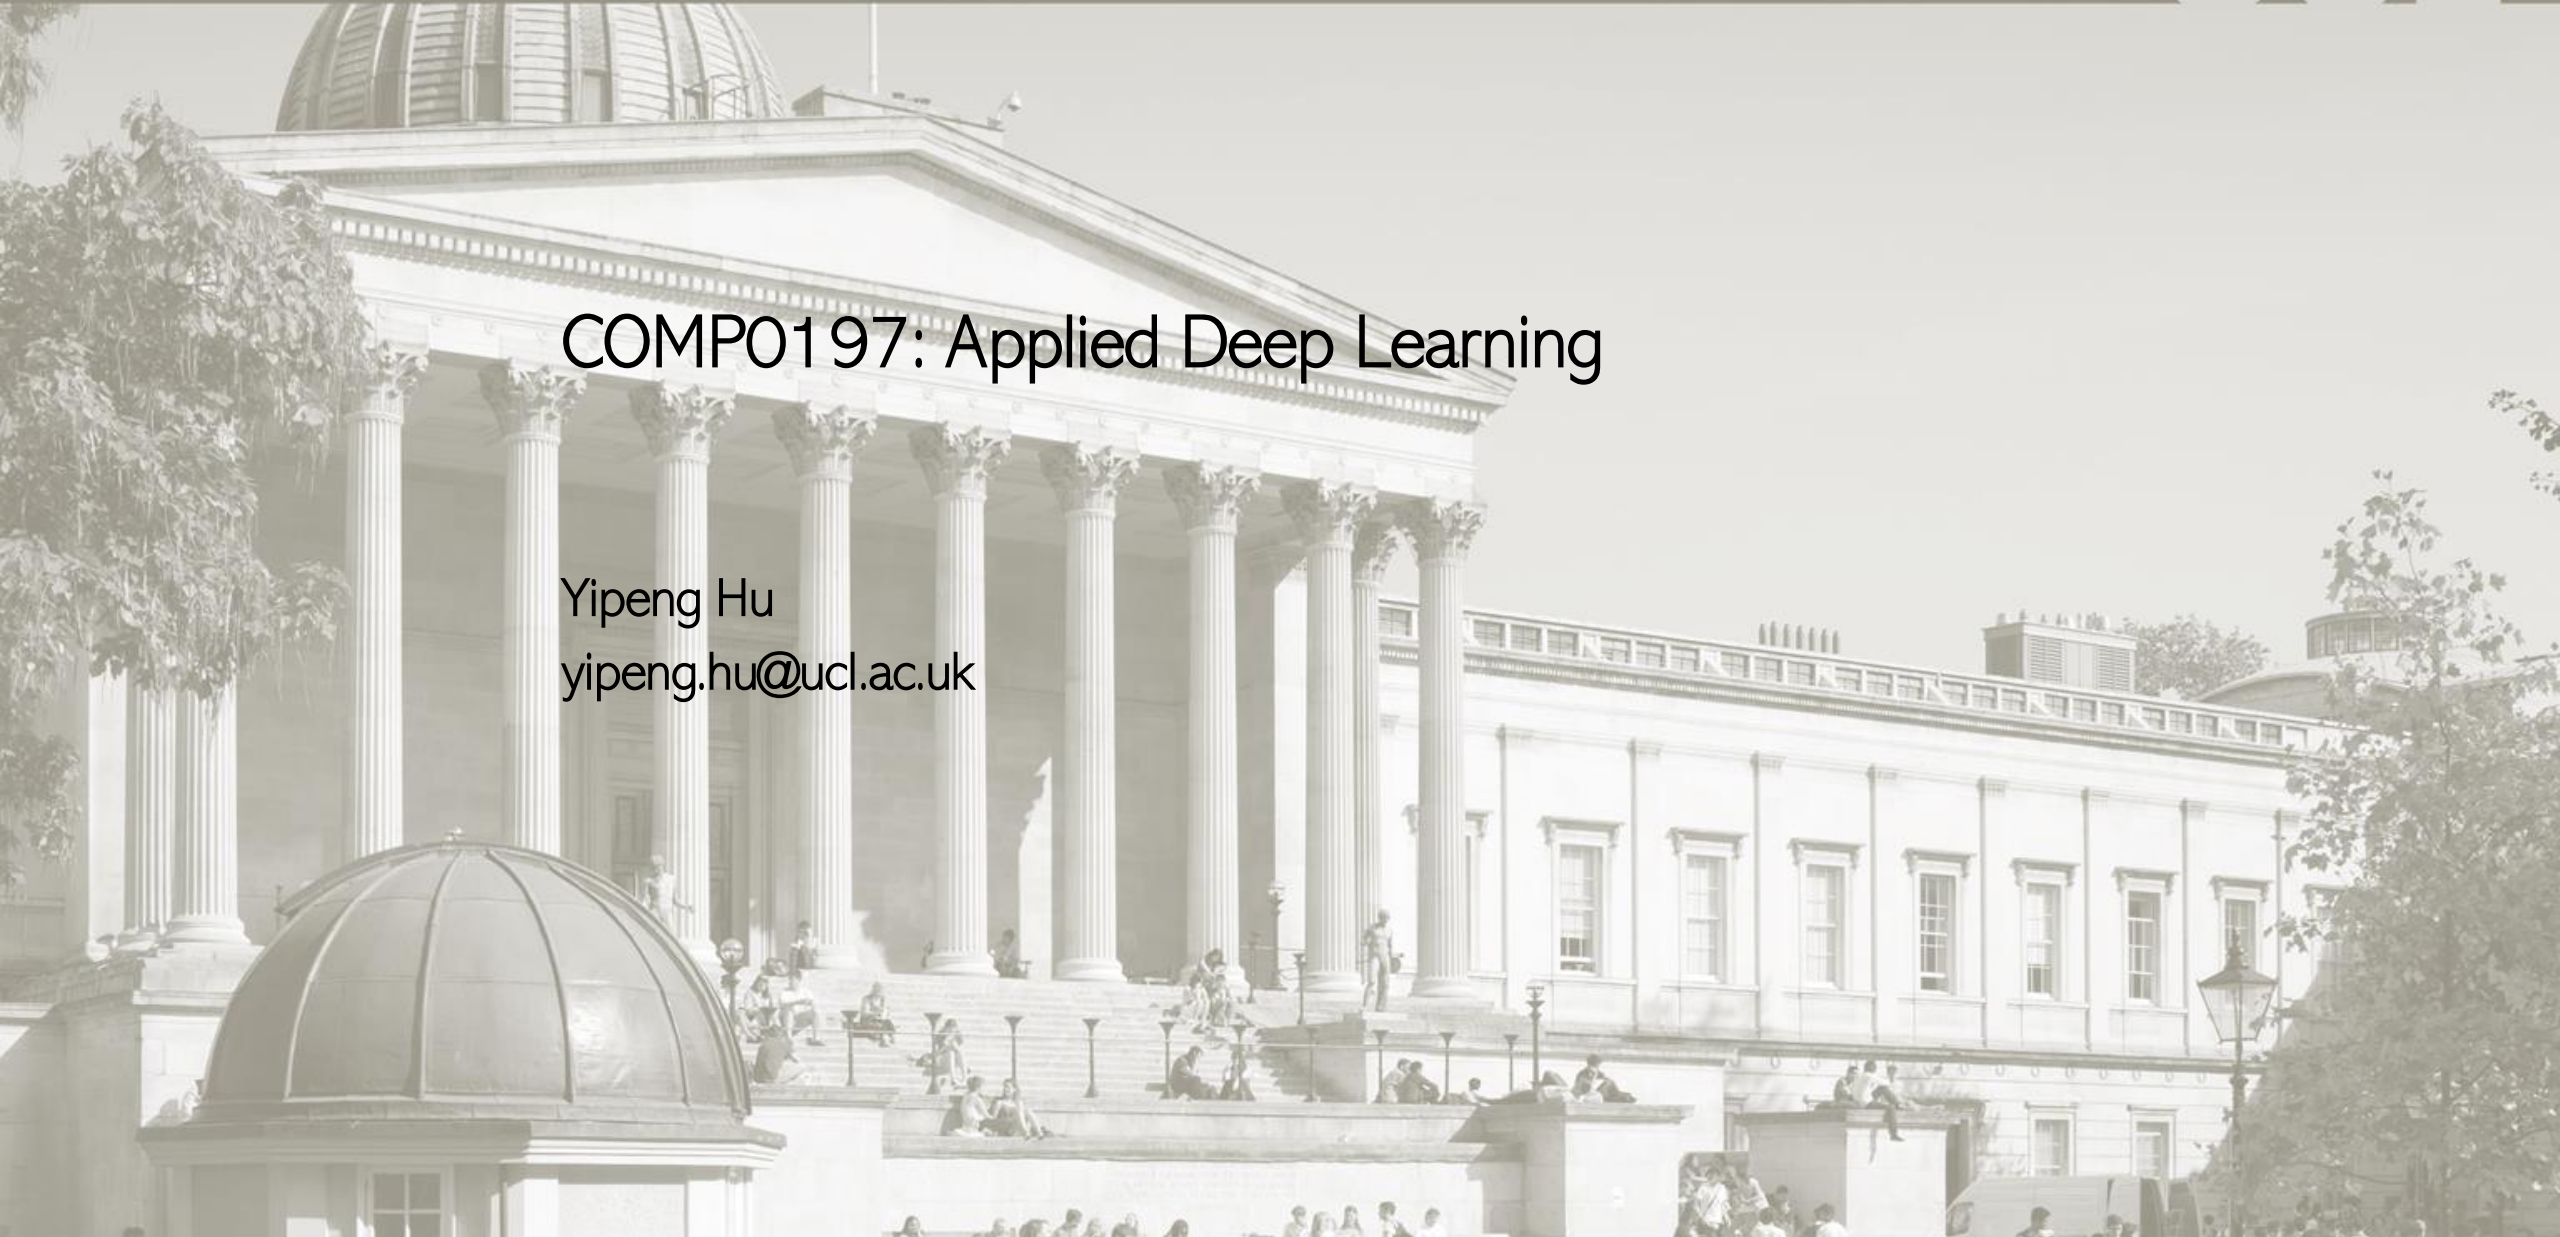


# COMP0197: Applied Deep Learning


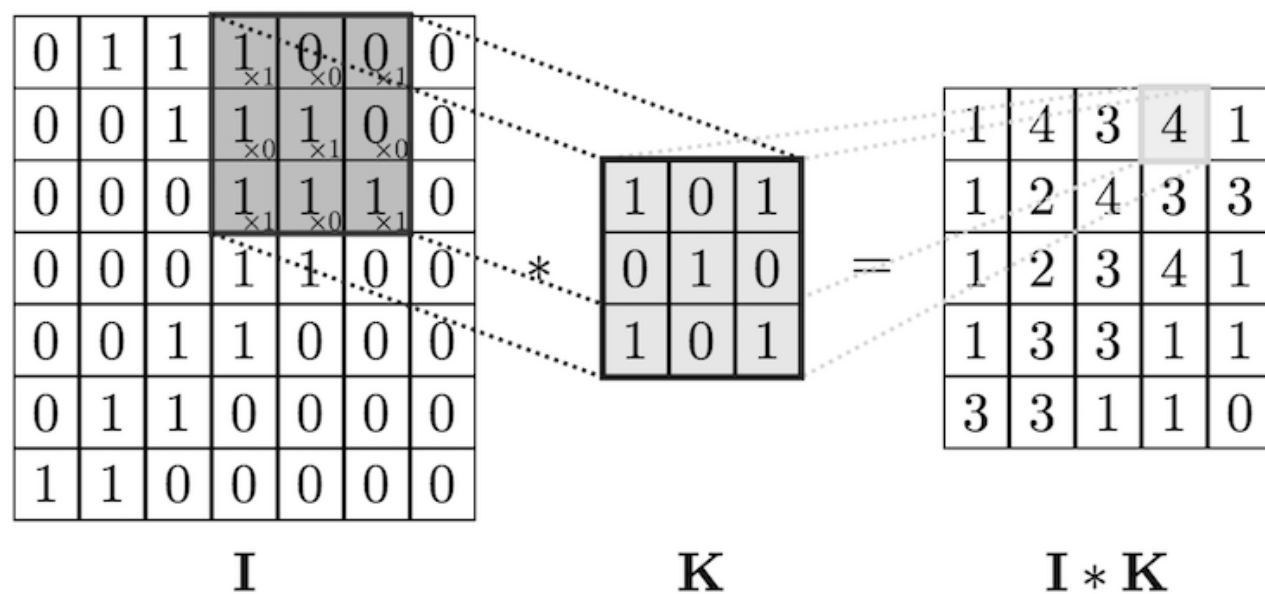
Yipeng Hu  
[yipeng.hu@ucl.ac.uk](mailto:yipeng.hu@ucl.ac.uk)



# Convolutional Neural Networks

## Discrete convolution

$$s(t) = (x * w)(t) = \sum_{a=-\infty}^{\infty} x(a)w(t-a)$$



$$\begin{bmatrix} x_1 & x_4 & x_7 \\ x_2 & x_5 & x_8 \\ x_3 & x_6 & x_9 \end{bmatrix} * \begin{bmatrix} k_1 & k_3 \\ k_2 & k_4 \end{bmatrix} = \begin{bmatrix} \sum \sum \begin{pmatrix} x_1 k_1 & x_4 k_3 \\ x_2 k_2 & x_5 k_4 \end{pmatrix} & \sum \sum \begin{pmatrix} x_4 k_1 & x_7 k_3 \\ x_5 k_2 & x_8 k_4 \end{pmatrix} \\ \sum \sum \begin{pmatrix} x_2 k_1 & x_5 k_3 \\ x_3 k_2 & x_6 k_4 \end{pmatrix} & \sum \sum \begin{pmatrix} x_5 k_1 & x_8 k_3 \\ x_6 k_2 & x_9 k_4 \end{pmatrix} \end{bmatrix}$$

$$= \begin{bmatrix} x_1 k_1 + x_2 k_2 + x_4 k_3 + x_5 k_4 & x_4 k_1 + x_5 k_2 + x_7 k_3 + x_8 k_4 \\ x_2 k_1 + x_3 k_2 + x_5 k_3 + x_6 k_4 & x_5 k_1 + x_6 k_2 + x_8 k_3 + x_9 k_4 \end{bmatrix}$$

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \end{bmatrix} = \begin{bmatrix} x_1 k_1 + x_2 k_2 + x_4 k_3 + x_5 k_4 \\ x_2 k_1 + x_3 k_2 + x_5 k_3 + x_6 k_4 \\ x_4 k_1 + x_5 k_2 + x_7 k_3 + x_8 k_4 \\ x_5 k_1 + x_6 k_2 + x_8 k_3 + x_9 k_4 \end{bmatrix}$$

$$\begin{bmatrix} x_1 & x_4 & x_7 \\ x_2 & x_5 & x_8 \\ x_3 & x_6 & x_9 \end{bmatrix} * \begin{bmatrix} k_1 & k_3 \\ k_2 & k_4 \end{bmatrix} = \begin{bmatrix} \sum \sum \begin{pmatrix} x_1 k_1 & x_4 k_3 \\ x_2 k_2 & x_5 k_4 \end{pmatrix} & \sum \sum \begin{pmatrix} x_4 k_1 & x_7 k_3 \\ x_5 k_2 & x_8 k_4 \end{pmatrix} \\ \sum \sum \begin{pmatrix} x_2 k_1 & x_5 k_3 \\ x_3 k_2 & x_6 k_4 \end{pmatrix} & \sum \sum \begin{pmatrix} x_5 k_1 & x_8 k_3 \\ x_6 k_2 & x_9 k_4 \end{pmatrix} \end{bmatrix}$$

$$= \begin{bmatrix} x_1 k_1 + x_2 k_2 + x_4 k_3 + x_5 k_4 & x_4 k_1 + x_5 k_2 + x_7 k_3 + x_8 k_4 \\ x_2 k_1 + x_3 k_2 + x_5 k_3 + x_6 k_4 & x_5 k_1 + x_6 k_2 + x_8 k_3 + x_9 k_4 \end{bmatrix}$$

$$\begin{bmatrix} k_1 & k_2 & 0 & k_3 & k_4 & 0 & 0 & 0 & 0 \\ 0 & k_1 & k_2 & 0 & k_3 & k_4 & 0 & 0 & 0 \\ 0 & 0 & 0 & k_1 & k_2 & 0 & k_3 & k_4 & 0 \\ 0 & 0 & 0 & 0 & k_1 & k_2 & 0 & k_3 & k_4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \end{bmatrix} = \begin{bmatrix} x_1 k_1 + x_2 k_2 + x_4 k_3 + x_5 k_4 \\ x_2 k_1 + x_3 k_2 + x_5 k_3 + x_6 k_4 \\ x_4 k_1 + x_5 k_2 + x_7 k_3 + x_8 k_4 \\ x_5 k_1 + x_6 k_2 + x_8 k_3 + x_9 k_4 \end{bmatrix}$$

## Convolution and cross-correlation

### Intuition

Translation-invariance, Parameter sharing, Sparse weights, Infinitely strong prior

### Sampling

Pooling, Strides, Interpolation, Un-pooling, Transpose convolution

## Anatomy of a convolutional layer

### Kernels

Dimensionality, receptive field, shape and size, kernel constraints, padding, depthwise separable convolution

### Convolution arithmetic

### Applied architectures

# Convolutional Neural Networks | Convolution and Cross-Correlation

## Convolution

















$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(m, n) K(i - m, j - n).$$

$$S(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i - m, j - n) K(m, n).$$

## Cross-correlation


$$S(i, j) = (I * K)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n).$$








## Convolutional Neural Networks | Intuition


## Translation invariance



## Translation invariance



## Translation invariance




## Parameter sharing

Regularisation effect\*, efficiency

$$\begin{bmatrix} w_1 & w_2 & 0 & w_3 & w_4 & 0 & 0 & 0 & 0 \\ 0 & w_1 & w_2 & 0 & w_3 & w_4 & 0 & 0 & 0 \\ 0 & 0 & 0 & w_1 & w_2 & 0 & w_3 & w_4 & 0 \\ 0 & 0 & 0 & 0 & w_1 & w_2 & 0 & w_3 & w_4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{bmatrix} = \begin{bmatrix} x_1 w_1 + x_2 w_2 + x_4 w_3 + x_5 w_4 \\ x_2 w_1 + x_3 w_2 + x_5 w_3 + x_6 w_4 \\ x_4 w_1 + x_5 w_2 + x_7 w_3 + x_8 w_4 \\ x_5 w_1 + x_6 w_2 + x_8 w_3 + x_9 w_4 \end{bmatrix}$$

$$\begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} & w_{1,5} & w_{1,6} & w_{1,7} & w_{1,8} & w_{1,9} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} & w_{2,5} & w_{2,6} & w_{2,7} & w_{2,8} & w_{2,9} \\ w_{3,1} & w_{3,2} & w_{3,3} & w_{3,4} & w_{3,5} & w_{3,6} & w_{3,7} & w_{3,8} & w_{3,9} \\ w_{4,1} & w_{4,2} & w_{4,3} & w_{4,4} & w_{4,5} & w_{4,6} & w_{4,7} & w_{4,8} & w_{4,9} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{bmatrix}$$




## Sparse weights

Memory, computation, statistical efficiency

$$\begin{bmatrix} w_1 & w_2 & 0 & w_3 & w_4 & 0 & 0 & 0 & 0 \\ 0 & w_1 & w_2 & 0 & w_3 & w_4 & 0 & 0 & 0 \\ 0 & 0 & 0 & w_1 & w_2 & 0 & w_3 & w_4 & 0 \\ 0 & 0 & 0 & 0 & w_1 & w_2 & 0 & w_3 & w_4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{bmatrix} = \begin{bmatrix} x_1 w_1 + x_2 w_2 + x_4 w_3 + x_5 w_4 \\ x_2 w_1 + x_3 w_2 + x_5 w_3 + x_6 w_4 \\ x_4 w_1 + x_5 w_2 + x_7 w_3 + x_8 w_4 \\ x_5 w_1 + x_6 w_2 + x_8 w_3 + x_9 w_4 \end{bmatrix}$$

$$\begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} & w_{1,5} & w_{1,6} & w_{1,7} & w_{1,8} & w_{1,9} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} & w_{2,5} & w_{2,6} & w_{2,7} & w_{2,8} & w_{2,9} \\ w_{3,1} & w_{3,2} & w_{3,3} & w_{3,4} & w_{3,5} & w_{3,6} & w_{3,7} & w_{3,8} & w_{3,9} \\ w_{4,1} & w_{4,2} & w_{4,3} & w_{4,4} & w_{4,5} & w_{4,6} & w_{4,7} & w_{4,8} & w_{4,9} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{bmatrix}$$




## Infinitely strong prior

Underfitting

$$\begin{bmatrix} w_1 & w_2 & 0 & w_3 & w_4 & 0 & 0 & 0 & 0 \\ 0 & w_1 & w_2 & 0 & w_3 & w_4 & 0 & 0 & 0 \\ 0 & 0 & 0 & w_1 & w_2 & 0 & w_3 & w_4 & 0 \\ 0 & 0 & 0 & 0 & w_1 & w_2 & 0 & w_3 & w_4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{bmatrix} = \begin{bmatrix} x_1 w_1 + x_2 w_2 + x_4 w_3 + x_5 w_4 \\ x_2 w_1 + x_3 w_2 + x_5 w_3 + x_6 w_4 \\ x_4 w_1 + x_5 w_2 + x_7 w_3 + x_8 w_4 \\ x_5 w_1 + x_6 w_2 + x_8 w_3 + x_9 w_4 \end{bmatrix}$$

$$\begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} & w_{1,5} & w_{1,6} & w_{1,7} & w_{1,8} & w_{1,9} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} & w_{2,5} & w_{2,6} & w_{2,7} & w_{2,8} & w_{2,9} \\ w_{3,1} & w_{3,2} & w_{3,3} & w_{3,4} & w_{3,5} & w_{3,6} & w_{3,7} & w_{3,8} & w_{3,9} \\ w_{4,1} & w_{4,2} & w_{4,3} & w_{4,4} & w_{4,5} & w_{4,6} & w_{4,7} & w_{4,8} & w_{4,9} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \\ x_8 \\ x_9 \end{bmatrix}$$





## Convolutional Neural Networks | Sampling

## Pooling

- Max pooling
- Average pooling


Invariance

No learnable parameters

Required for down-sampling?



Conv weight  $\rightarrow$  activation  $\rightarrow$  pooling?

Moving windows with strides




## Strides for down-sampling

- Strides  $> 1$
- Pooling
- Convolution with strides?




## Interpolation for down-sampling and up-sampling

Bilinear, bicubic, spline.



## Interpolation for down-sampling and up-sampling

Is bilinear interpolation a linear operation?



## “Un-pooling” for up-sampling

e.g. a fixed position for the max  
and zeros everywhere else

0.1	0.5	<b>1.2</b>	-0.7
<b>0.8</b>	-0.2	-0.5	0.3
0.4	<b>0.9</b>	-0.1	-0.2
-0.6	0.1	<b>0.5</b>	0.3

max-pooling

0.8	1.2
0.9	0.5

		x	
x			
	x		
		x	

max locations

0	0	<b>0.5</b>	0
<b>1.3</b>	0	0	0
0	<b>0.4</b>	0	0
0	0	<b>0.1</b>	0

unpooling

1.3	0.5
0.4	0.1

## Transpose convolution

- Fractionally strided convolution
- Deconvolution
- Up-sampling with parameters

$$\mathbf{y}^{(N \times 1)} = \mathbf{W}^{(N \times M)} \mathbf{x}^{(M \times 1)}$$

$$\mathbf{y} = \mathbf{W}\mathbf{x}$$

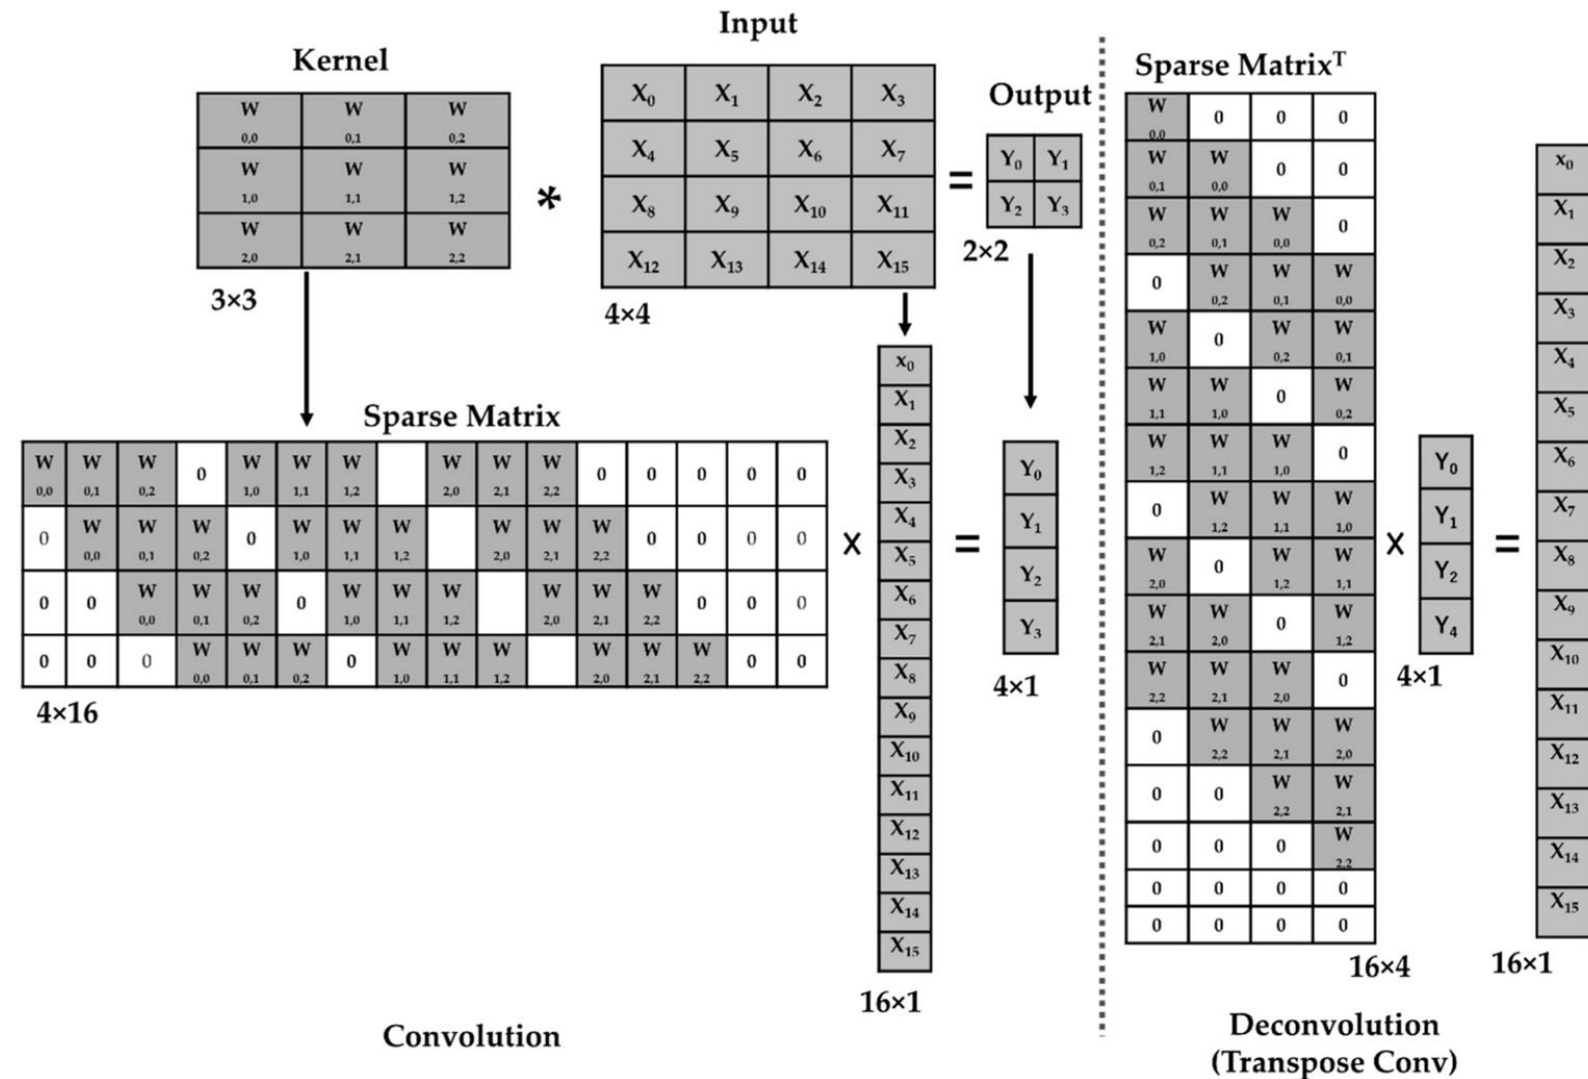
Diagonal-constant Toeplitz matrix

$$W_{i,j} = \text{Const}_{i,j}$$




$$\hat{W}_{j,i} = \text{Const}_{i,j}$$

The reverse of convolution

$$\hat{\mathbf{W}}^T \mathbf{y} = \mathbf{x}$$



## Transpose convolution

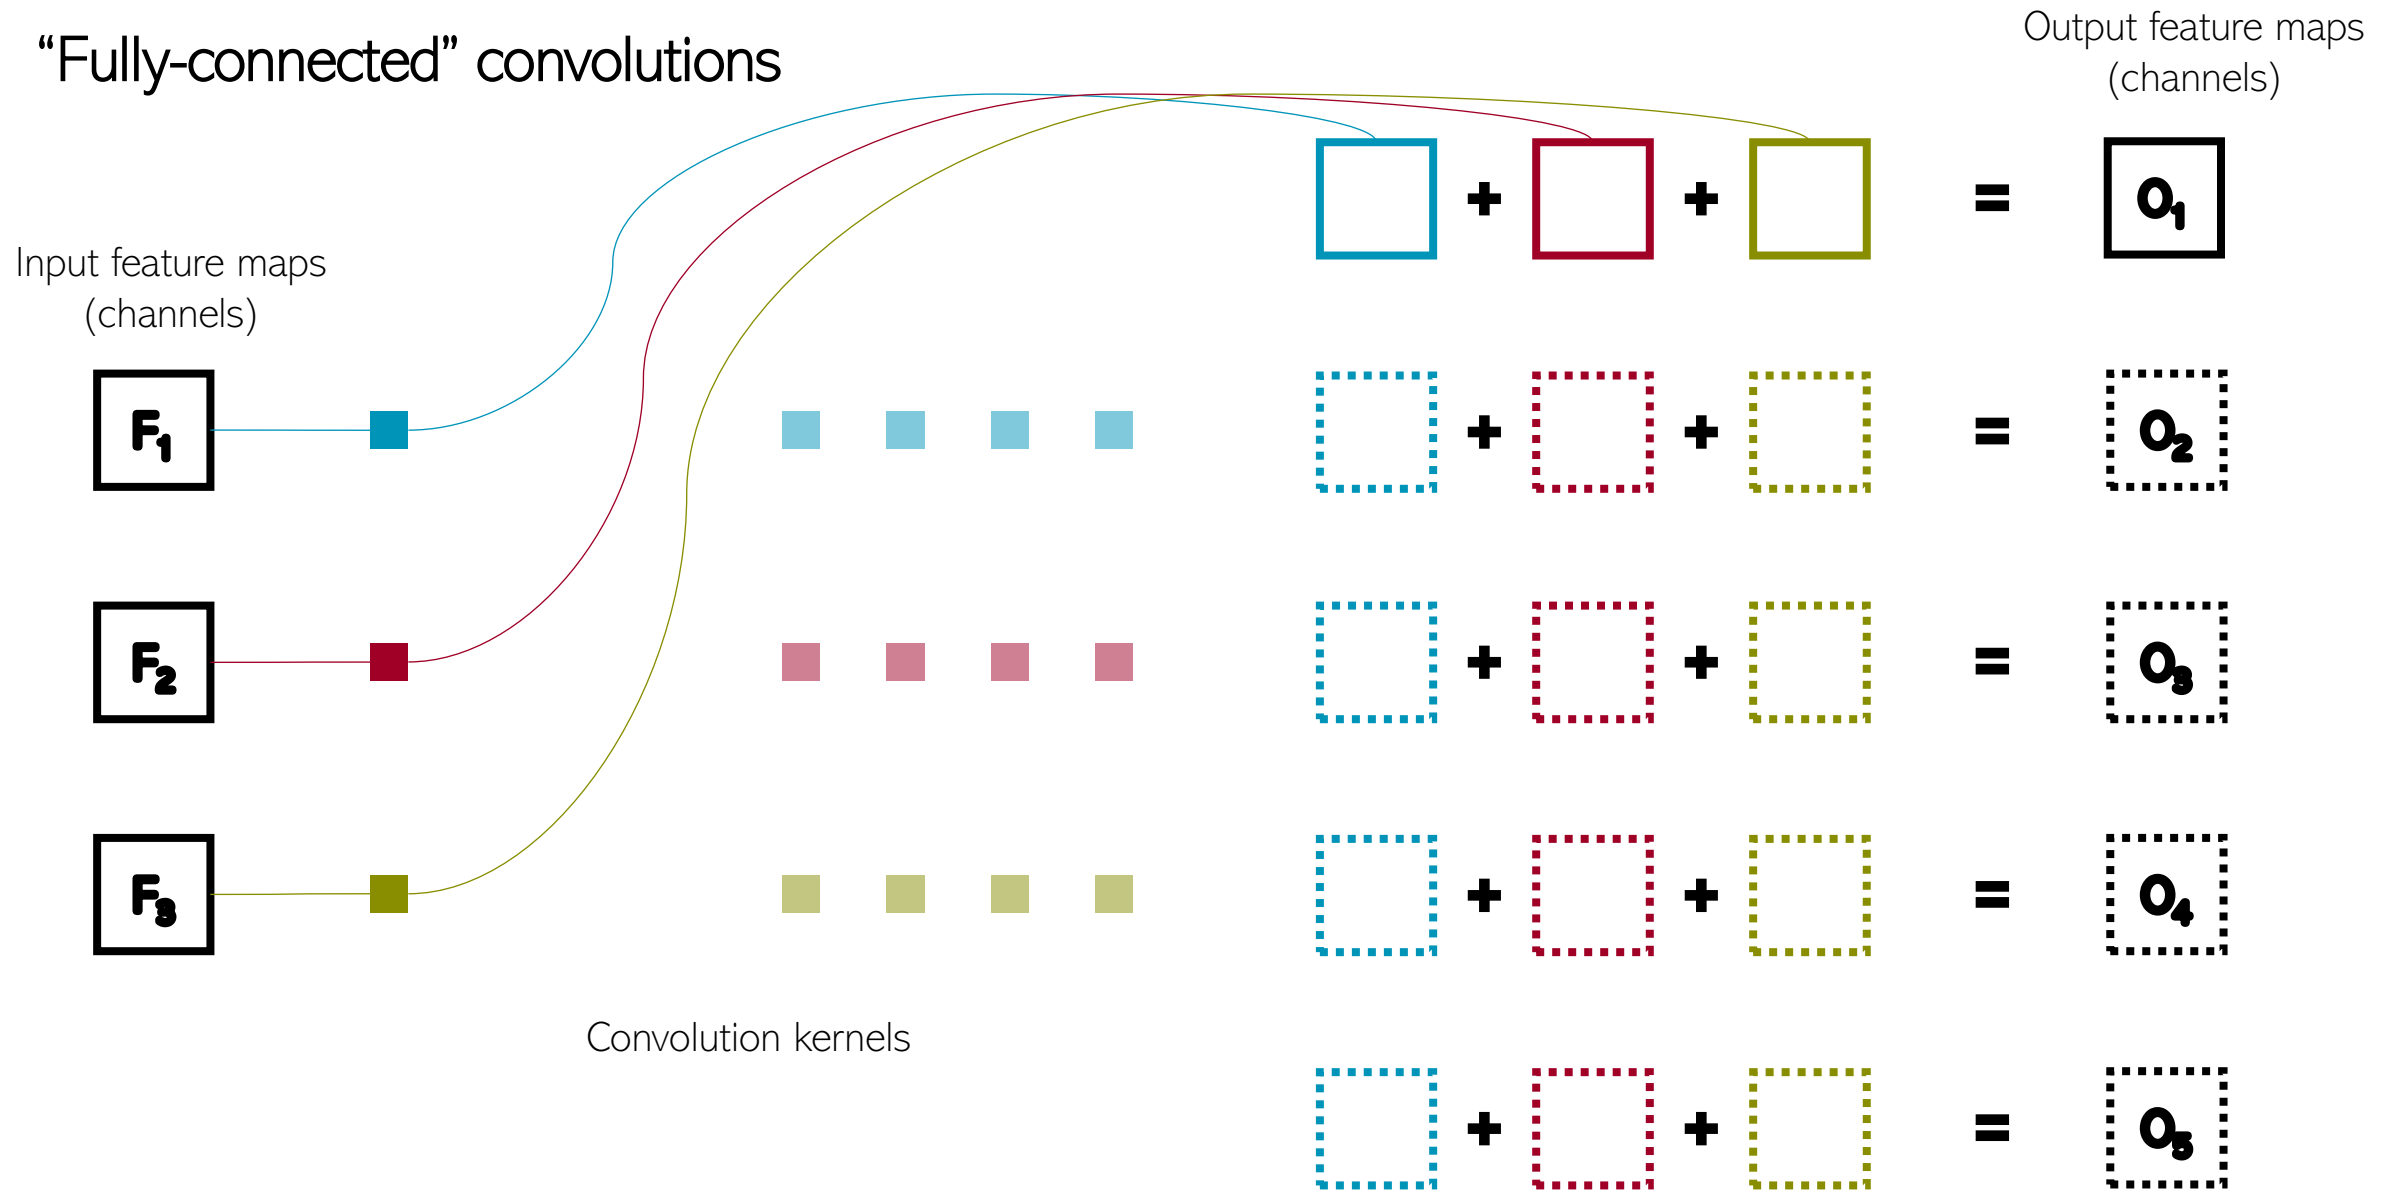


Transpose convolution vs. conv + interpolation, for up-sampling?



## Convolutional Neural Networks | Anatomy of a Convolutional Layer

## “Fully-connected” convolutions



## Tensors

Input/output feature maps

Tensor for 2D feature maps

$N \times H \times W \times C$  (TensorFlow)

$N \times C \times H \times W$  (PyTorch)

N: no. of data in a minibatch (batch size)

H: height of feature map

W: width of feature map

Tensor for 1D feature vector

$N \times L \times C$  (TensorFlow)

$N \times C \times L$  (PyTorch)


L: length of feature vector

Tensor for 3D feature tensor

$N \times D \times H \times W \times C$  (TensorFlow)

$N \times C \times D \times H \times W$  (PyTorch)




D: depth of feature tensor








## Convolutional Neural Networks | Kernels

## Dimensionality




1D



2D





3D





## Receptive field

neuroscience / physics





## Receptive field

Fully-connected layers vs. convolutional layers vs. convolution with strides\*



## Receptive field

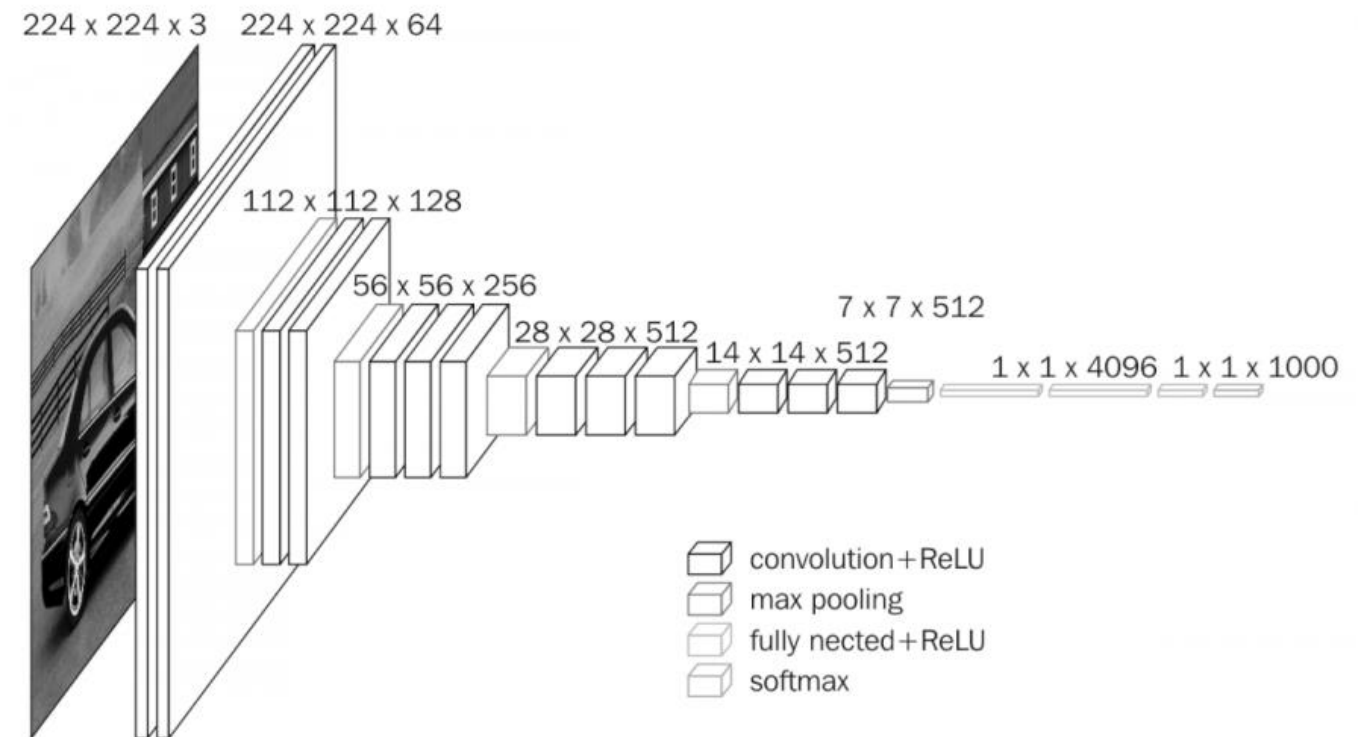
Dilated kernels (Atrous convolution)





## Shape and size

- Dilated
- Adaptive\*
- Lower-dimensional 2.5D kernels\*
- VGG: 3x3 kernels
- Bottleneck 1x1 kernels vs. fully-connected layers



Published as a conference paper at ICLR 2015


## VERY DEEP CONVOLUTIONAL NETWORKS FOR LARGE-SCALE IMAGE RECOGNITION

**Karen Simonyan\*** & **Andrew Zisserman\***

Visual Geometry Group, Department of Engineering Science, University of Oxford  
{karen, az}@robots.ox.ac.uk

## Shape and size

- 1x1 kernels
- “Channel-wise” linear projection (fully-connected layer)
- Dimensionality reduction for information “bottleneck” layer
- Upsampling







## Kernel constraints

- Soft constraints: Regularise, e.g.  $L^2$ -norm\*
- Hard constraint: MaxNorm\*

## Padding

No padding, arbitrary padding, half padding, full padding

- No strides
- With strides\*
- Transpose\*
- Dilated\*



## Depthwise separable kernels

- Separable convolution

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * [1 \quad 0 \quad -1]$$

$$\mathbf{K} = \mathbf{u} * \mathbf{v}$$

$$\mathbf{I} * \mathbf{K} = \mathbf{I} * \mathbf{u} * \mathbf{v} \text{ (associativity)}$$

Computational complexity

$$O(M \times N \times 3 \times 3) \rightarrow O(M \times N \times (3 + 3))$$

$$O(M \times N \times m \times n) \rightarrow O(M \times N \times (m + n))$$

Singular value decomposition

$$\mathbf{K} = [\mathbf{u}_1 \quad \mathbf{u}_2 \quad \mathbf{u}_3] \begin{bmatrix} S_1 & 0 & 0 \\ 0 & S_2 & 0 \\ 0 & 0 & S_3 \end{bmatrix} [\mathbf{v}_1 \quad \mathbf{v}_2 \quad \mathbf{v}_3]^T$$

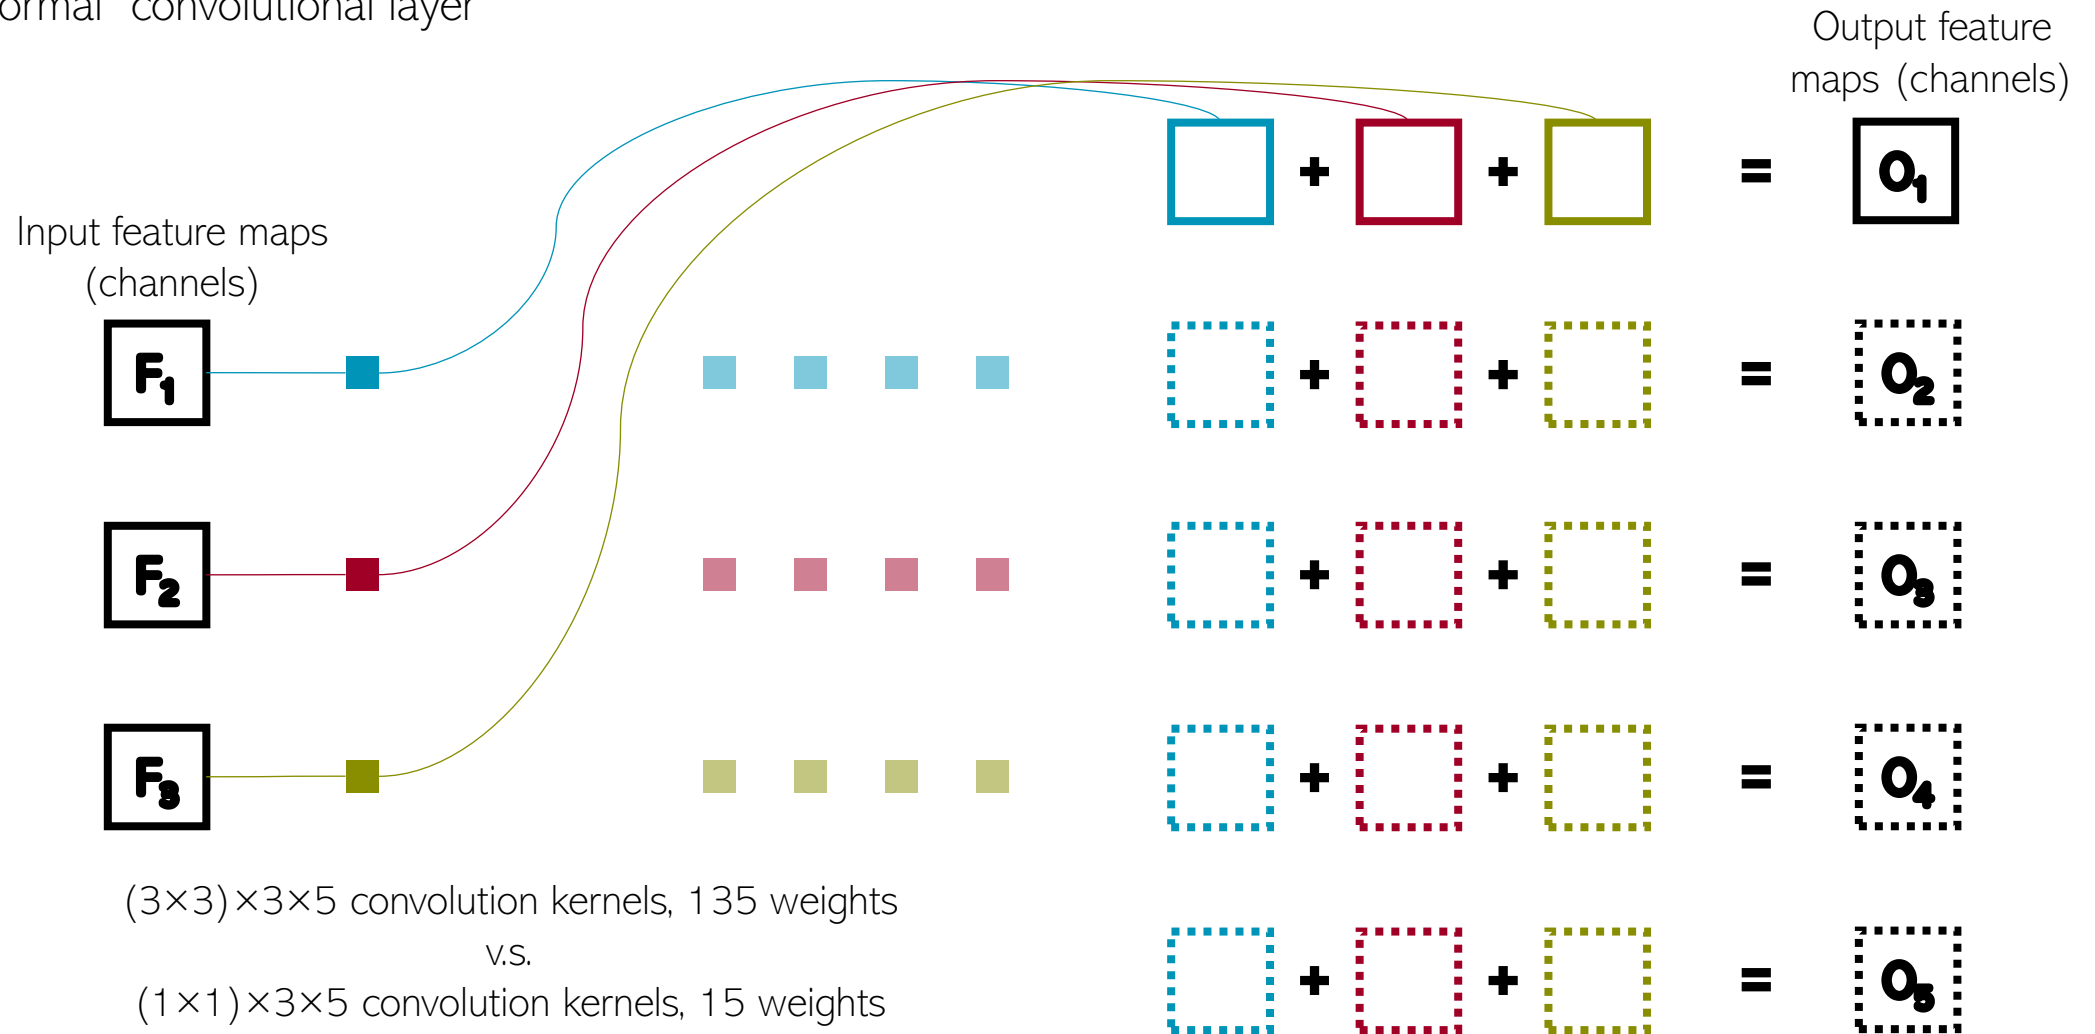
If  $\text{rank}(\mathbf{K}) = 1$  (test the number of non-singular values / linearly-independent vectors)

$$\mathbf{K} = S_1 \mathbf{u}_1 \mathbf{v}_1^T = S_1 \mathbf{u}_1 * \mathbf{v}_1^T \text{ (definition of convolution)}$$

Reduced degrees of freedom from 9 to 6

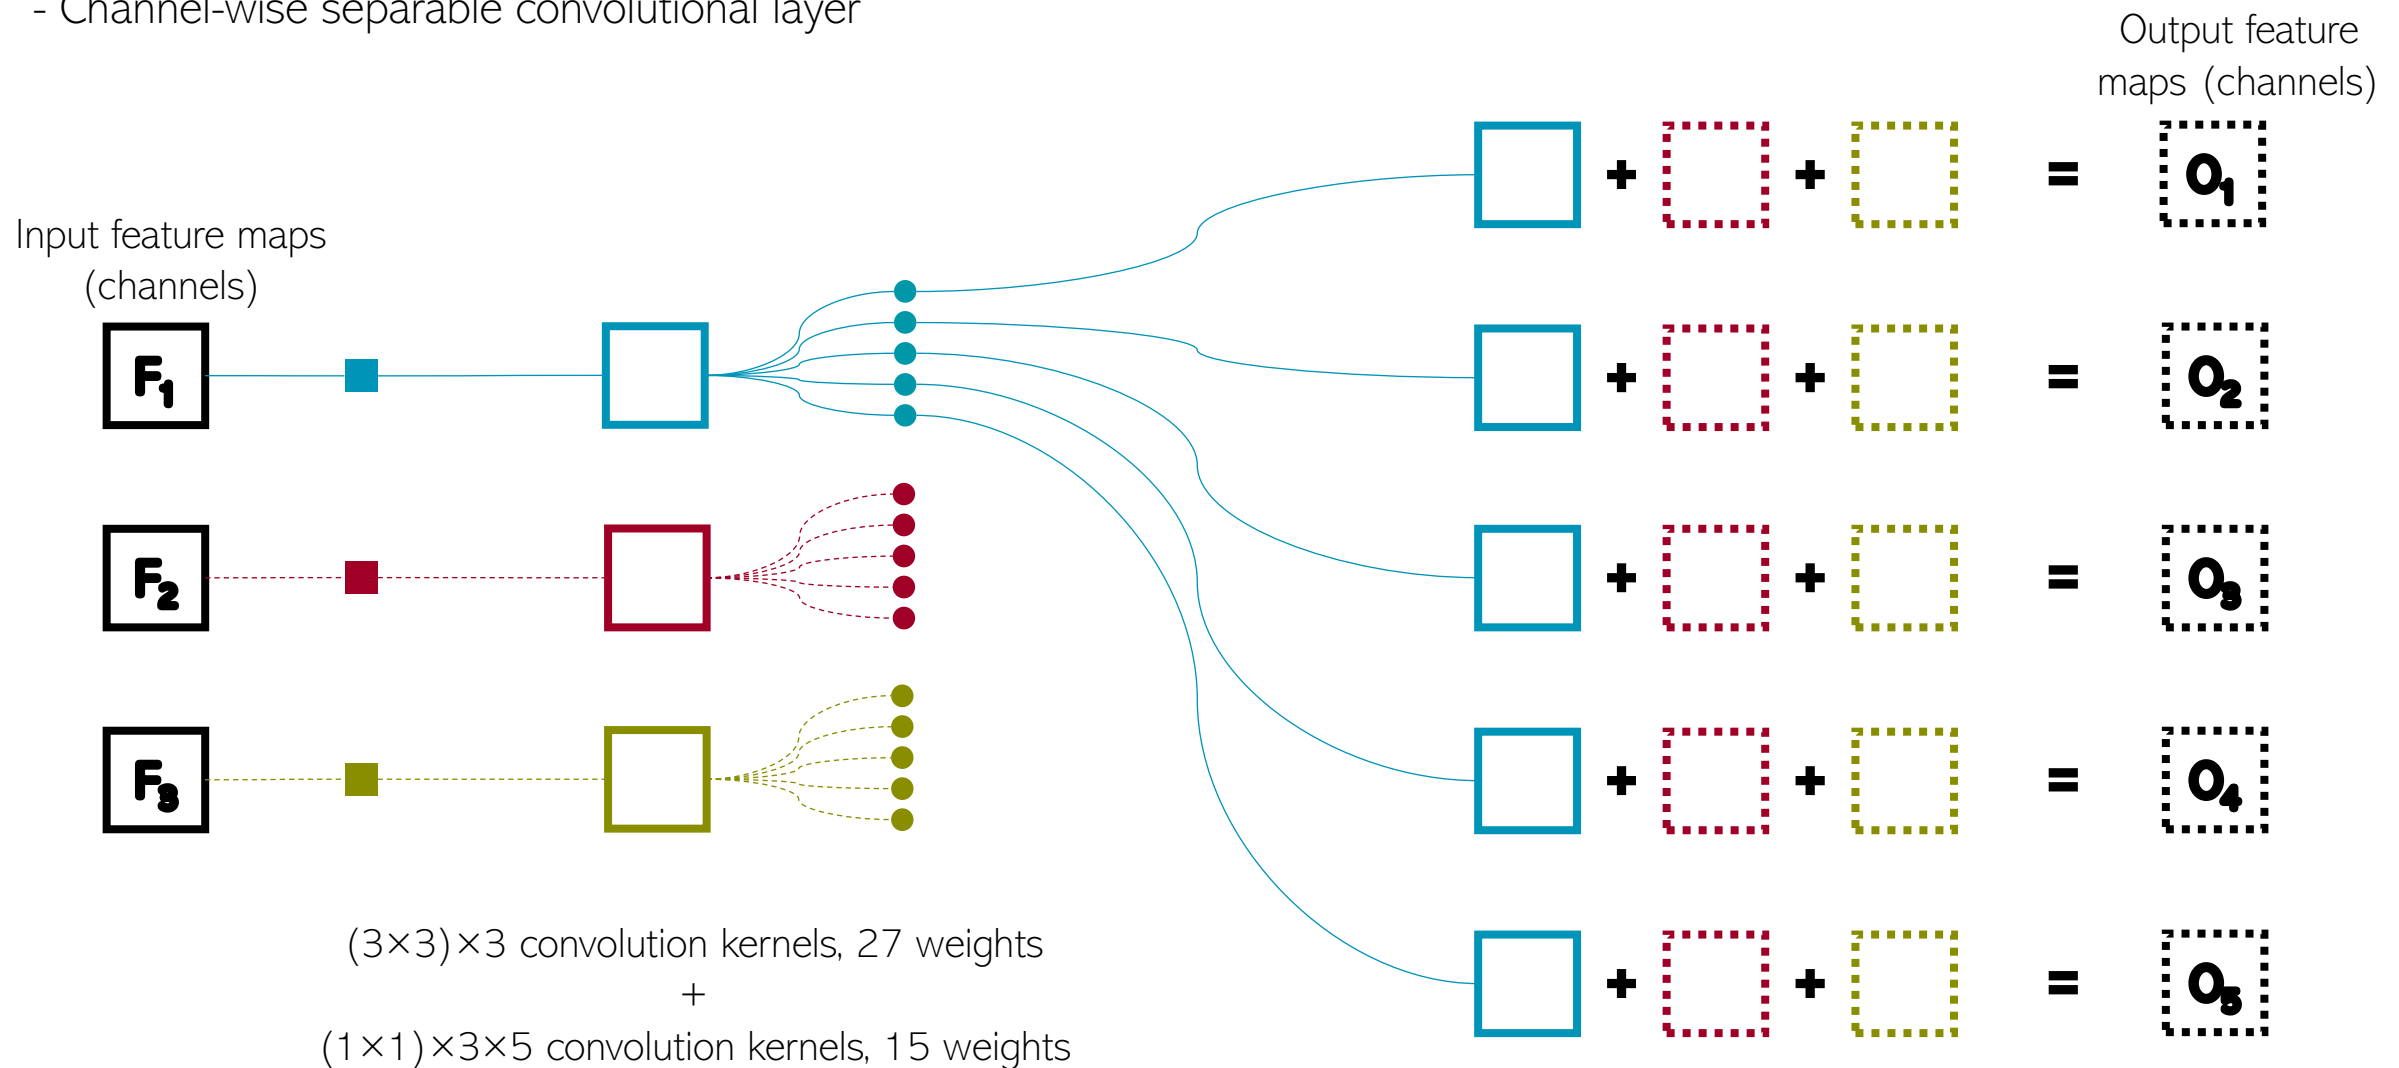
## Depthwise separable kernels

- Revisit the “normal” convolutional layer



## Depthwise separable kernels

- Channel-wise separable convolutional layer



## Depthwise separable kernels

- $(3 \times 3) \times 3 \times 5 = 135$  vs.  $(3 \times 3) \times 3 + (1 \times 1) \times 3 \times 5 = 42$
- $(5 \times 5) \times 64 \times 128 = 204,800$  vs.  $(5 \times 5) \times 64 + (1 \times 1) \times 64 \times 128 = 9,792$
- Model scaling
  - \* Width – number of (input and output) channels
  - \* Depth – number of layers
  - \* Resolution – feature map size
- “MobileNet”, scaling vs. performance
- “EfficientNet”, neural architecture search



## Convolutional Neural Networks | Convolution Arithmetic

Dumoulin, V. and Visin, F., 2016.

**A guide to convolution arithmetic for deep learning.**

arXiv preprint arXiv:1603.07285.


[https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

Things need to consider

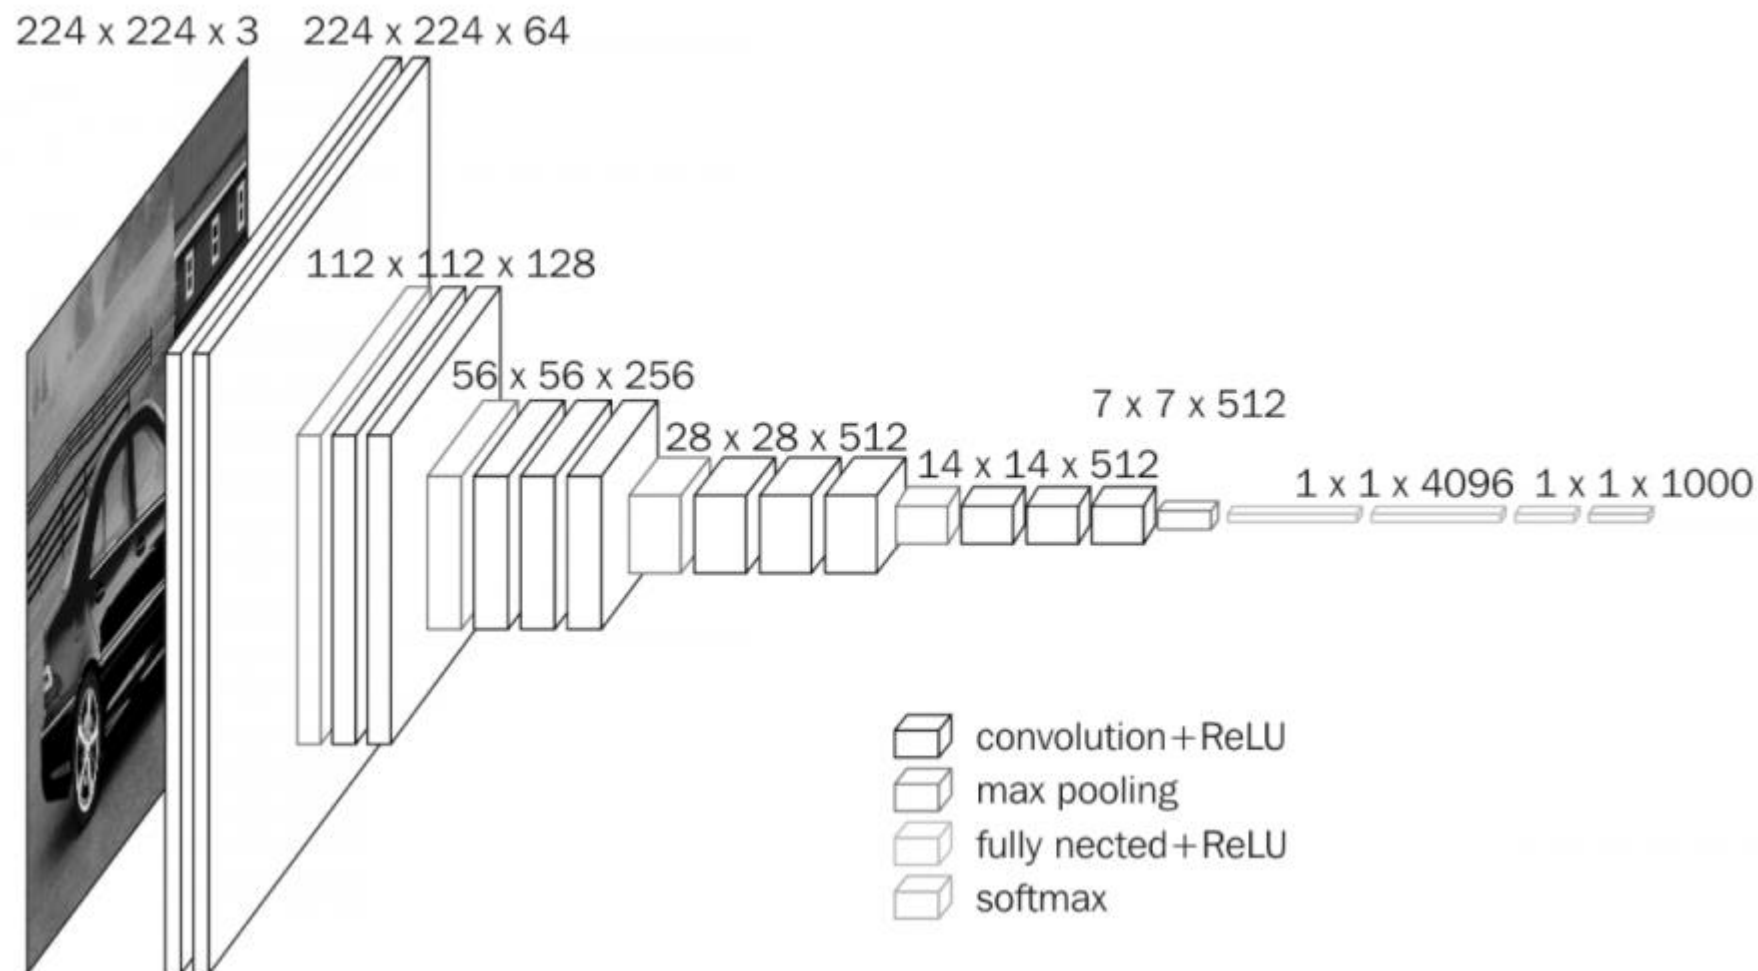
- Input/output feature map size and even/odd
- Kernel size and even/odd
- Strides
- Padding
- Kernel dilation rate
- Transpose convolution
- Pooling

## Convolutional Neural Networks | Applied Architectures


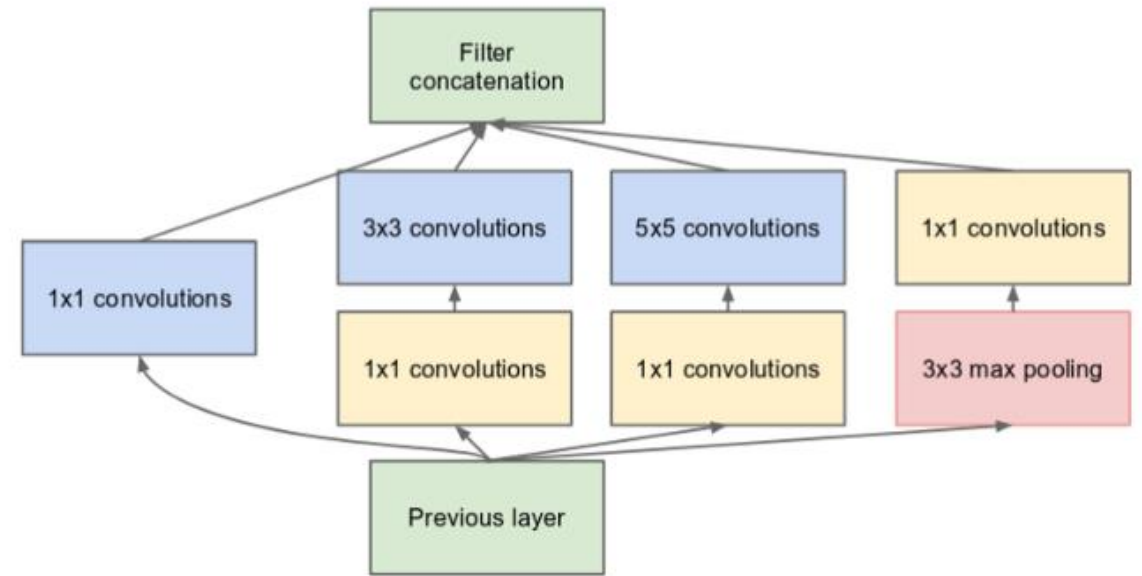
## AlexNet




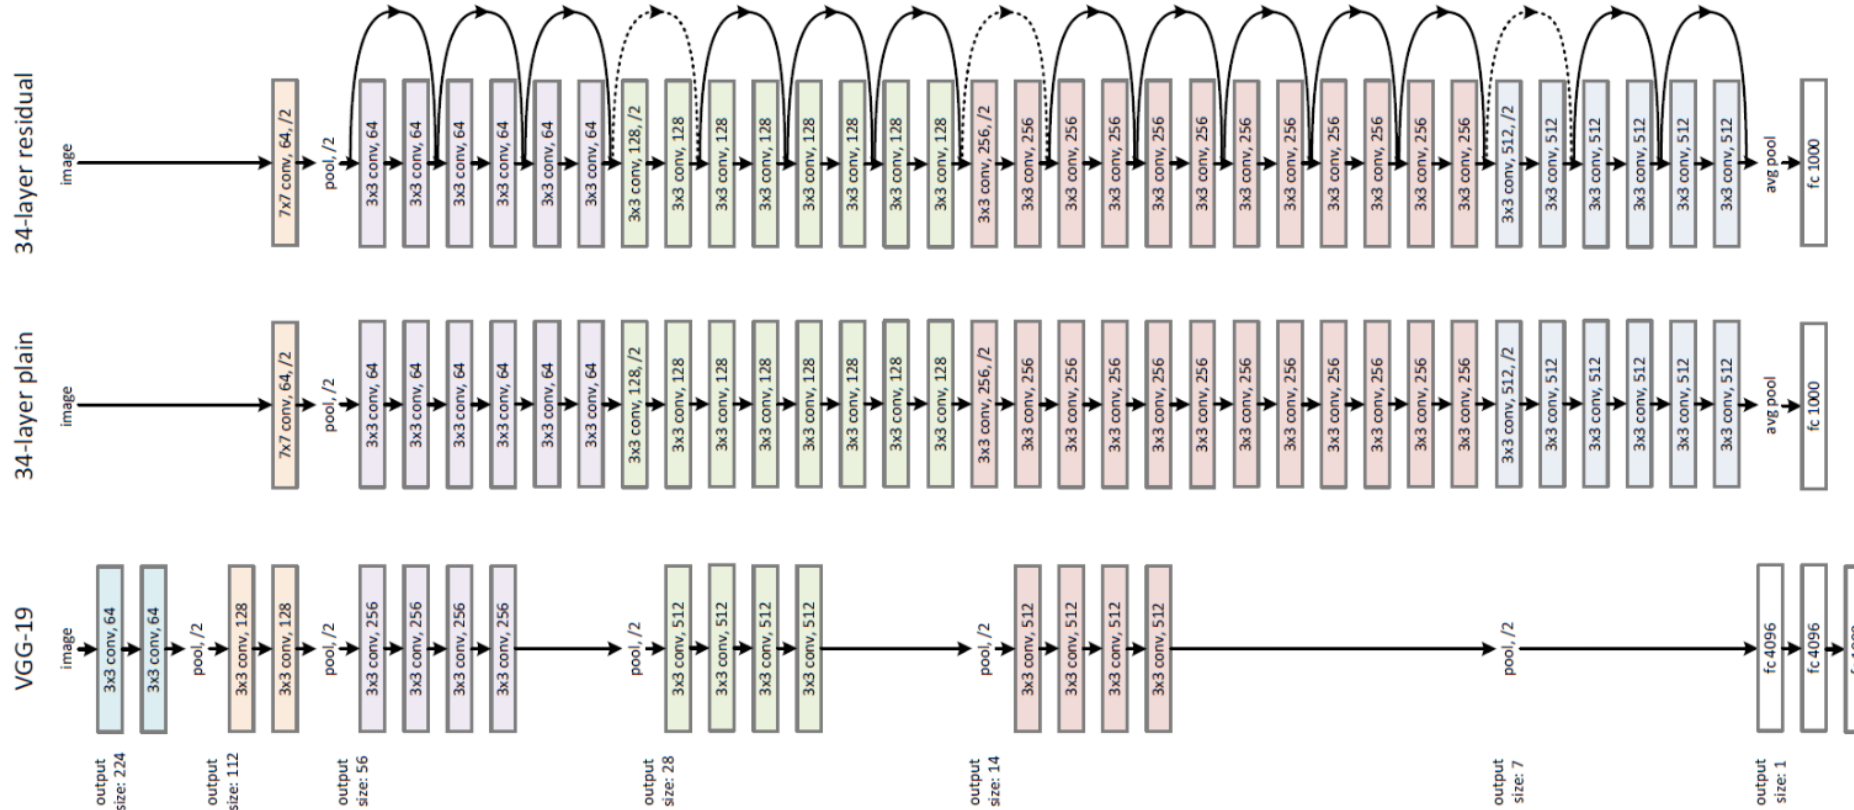
## VGG




## Inception



## ResNet

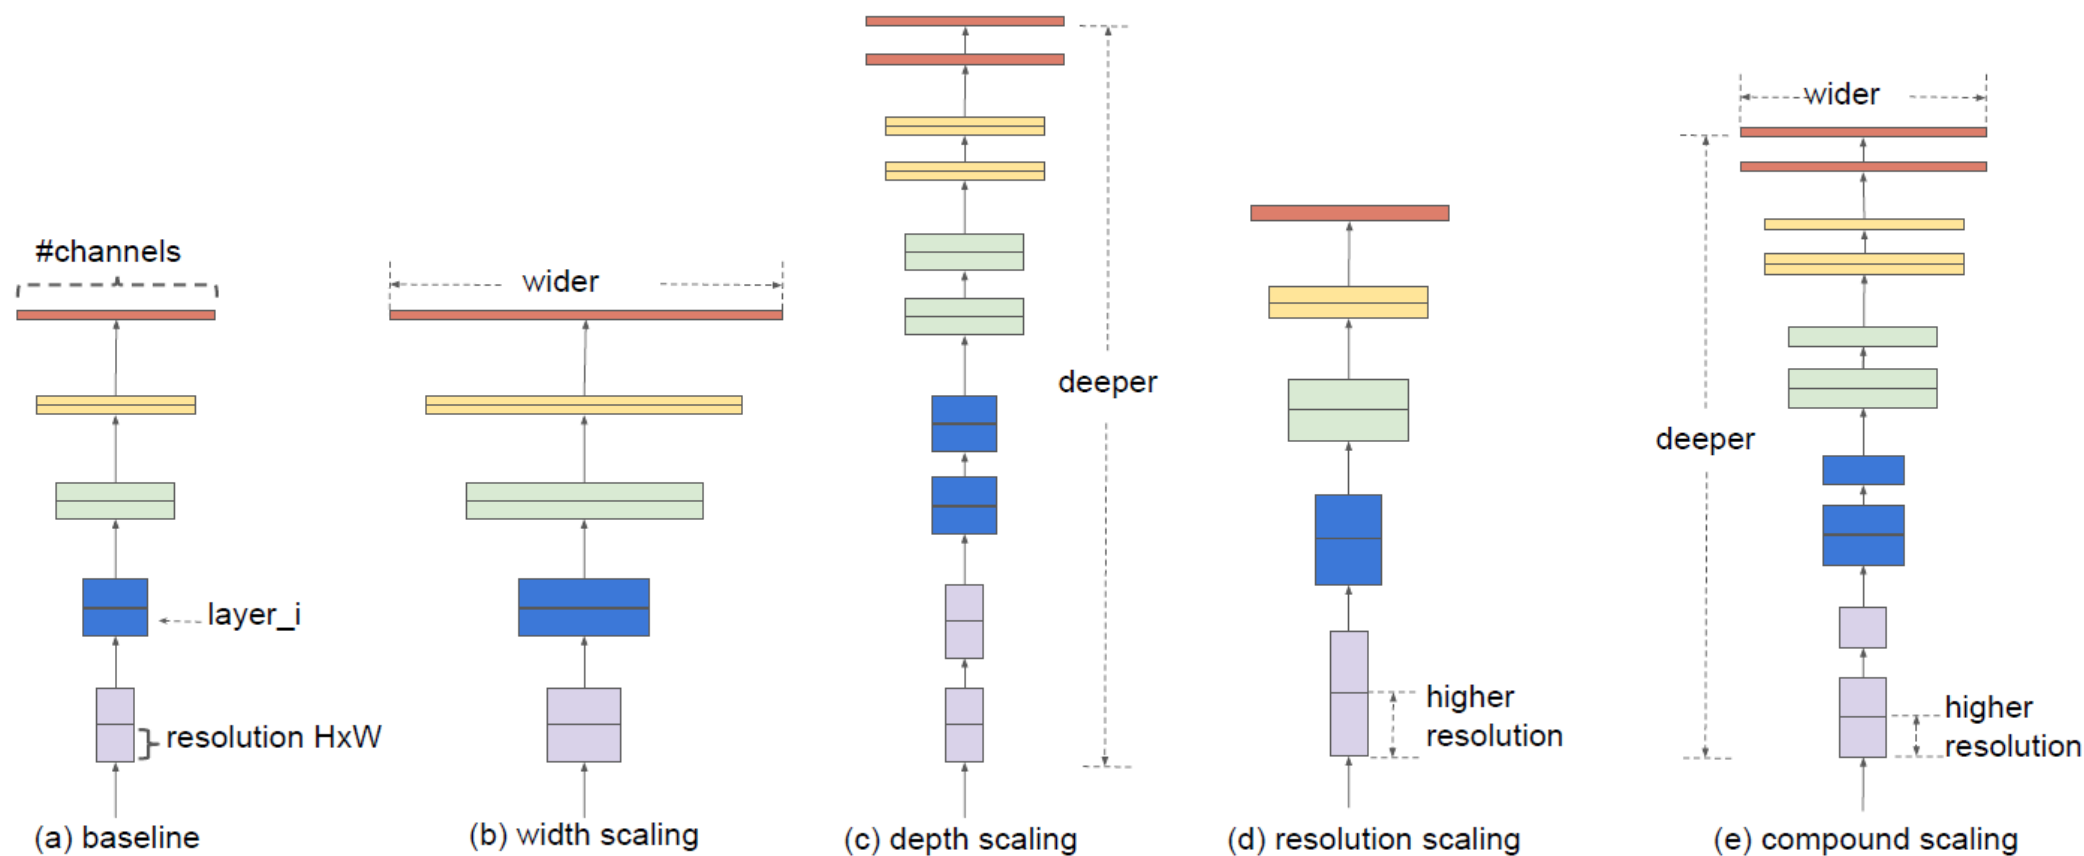


## DenseNet






## EfficientNet



## UNet



## Convolution and cross-correlation

### Intuition

Translation-invariance, Parameter sharing, Sparse weights, Infinitely strong prior

### Sampling

Pooling, Strides, Interpolation, Un-pooling, Transpose convolution


## Anatomy of a convolutional layer

### Kernels

Dimensionality, receptive field, shape and size, kernel constraints, padding, depthwise separable convolution

### Convolution arithmetic

### Applied architectures



Using a different CNN for the “image classification” tutorial  
Adding/reducing a resolution level in the “image segmentation” tutorial