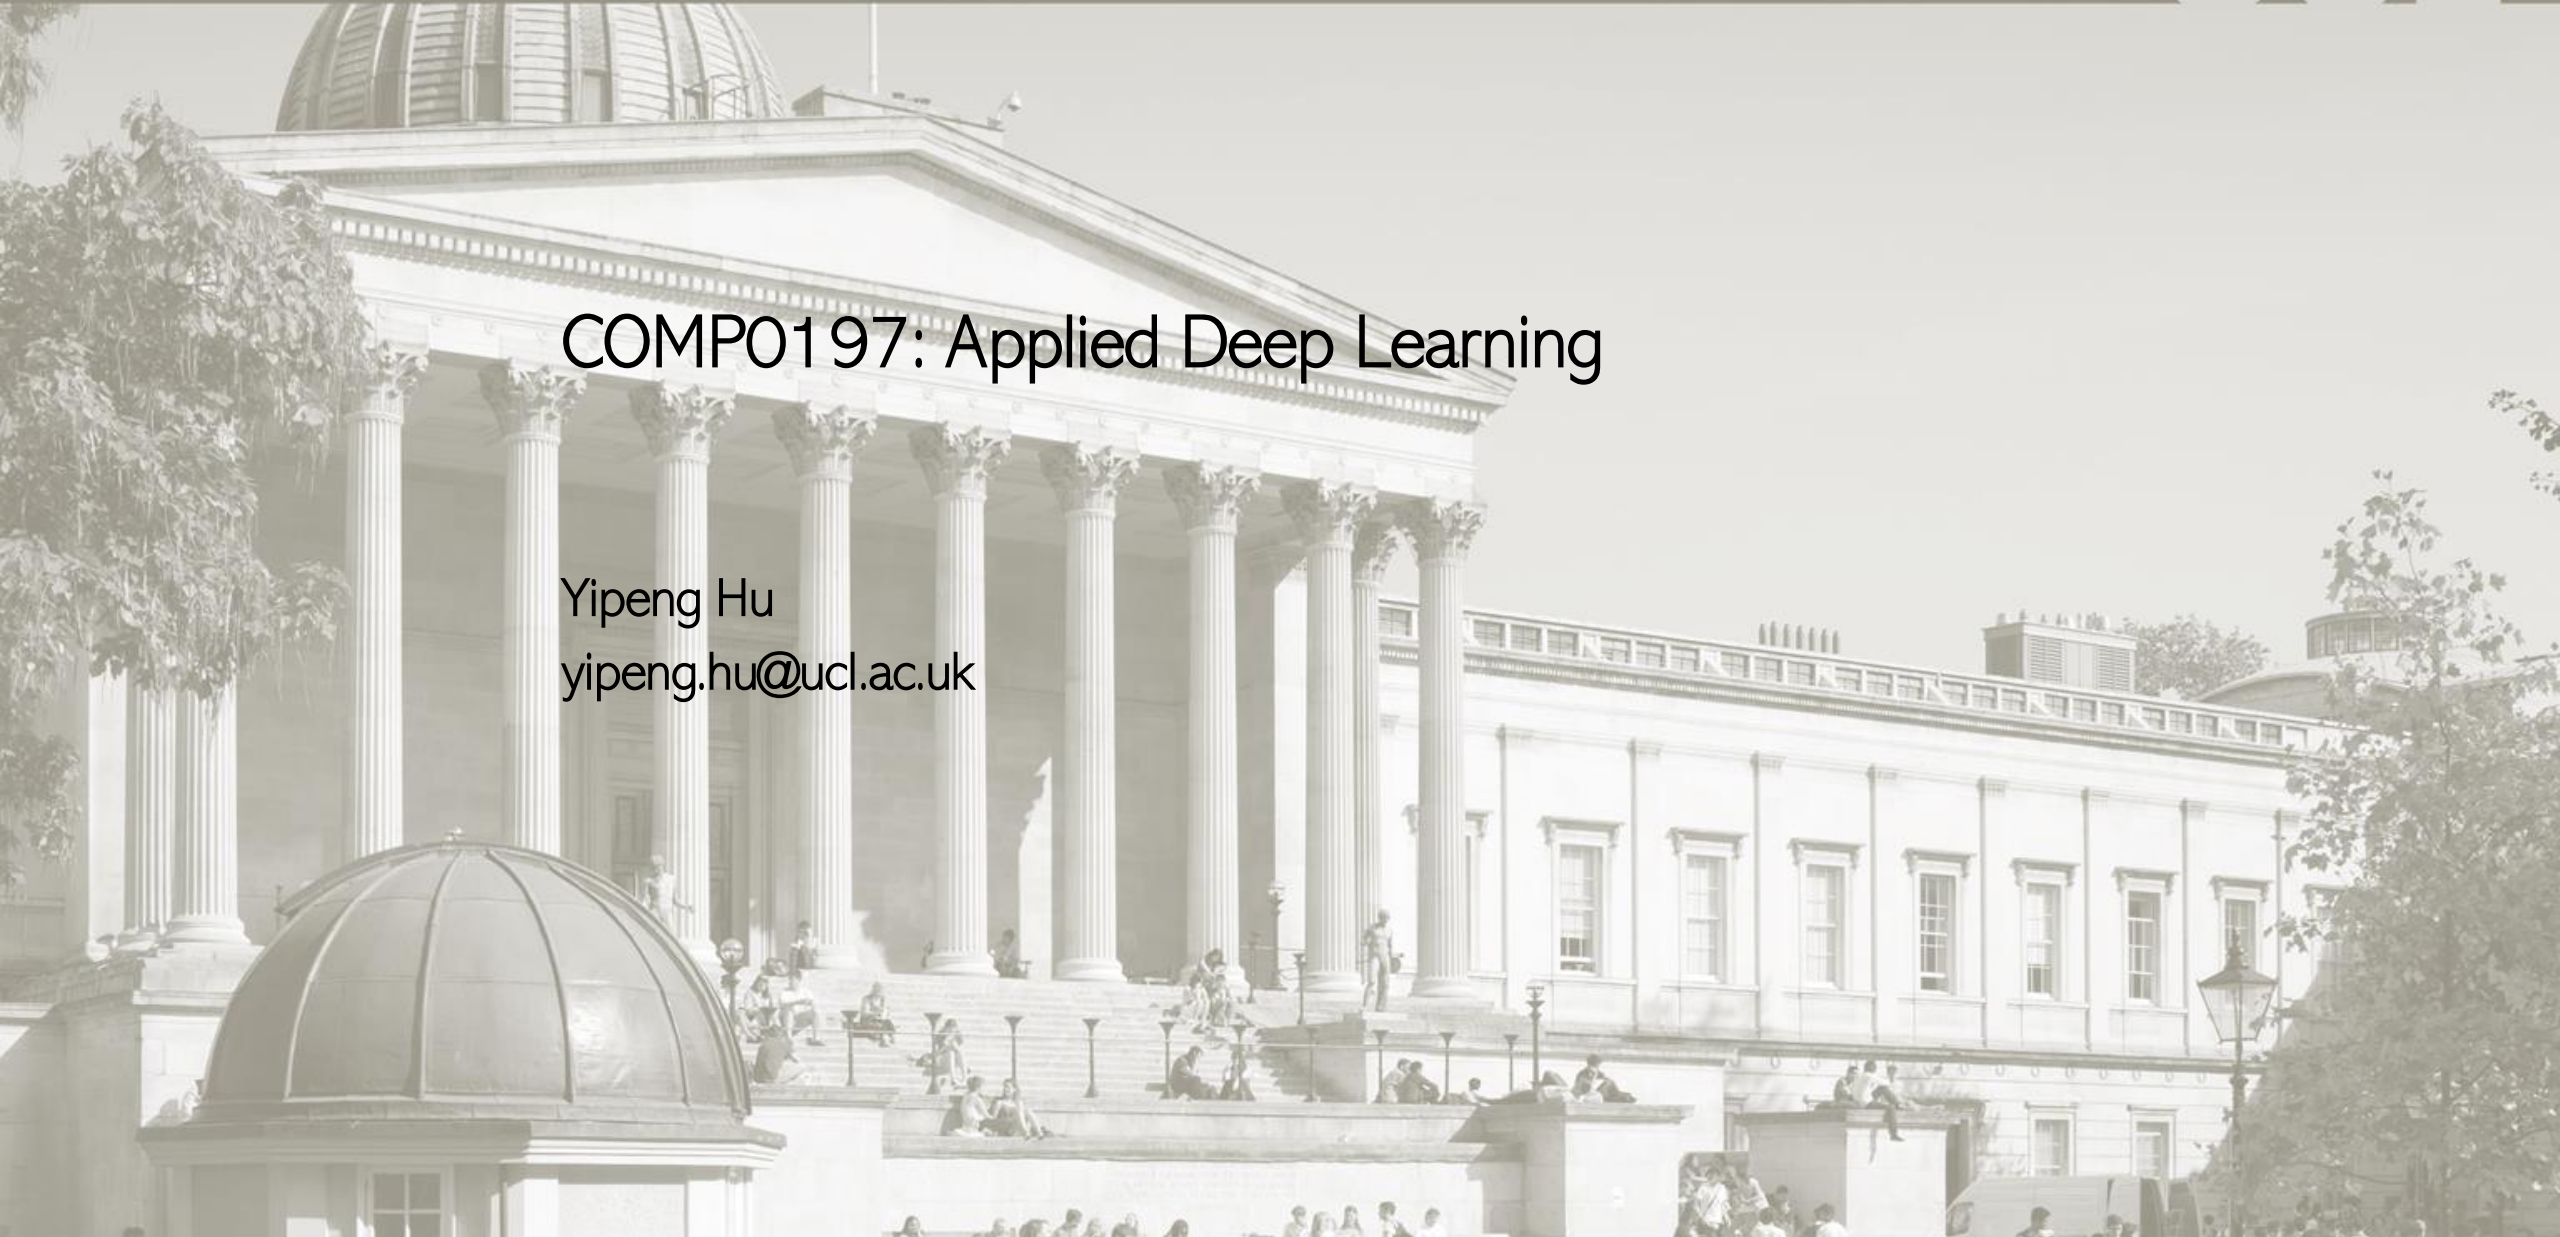


COMP0197: Applied Deep Learning

Yipeng Hu
yipeng.hu@ucl.ac.uk



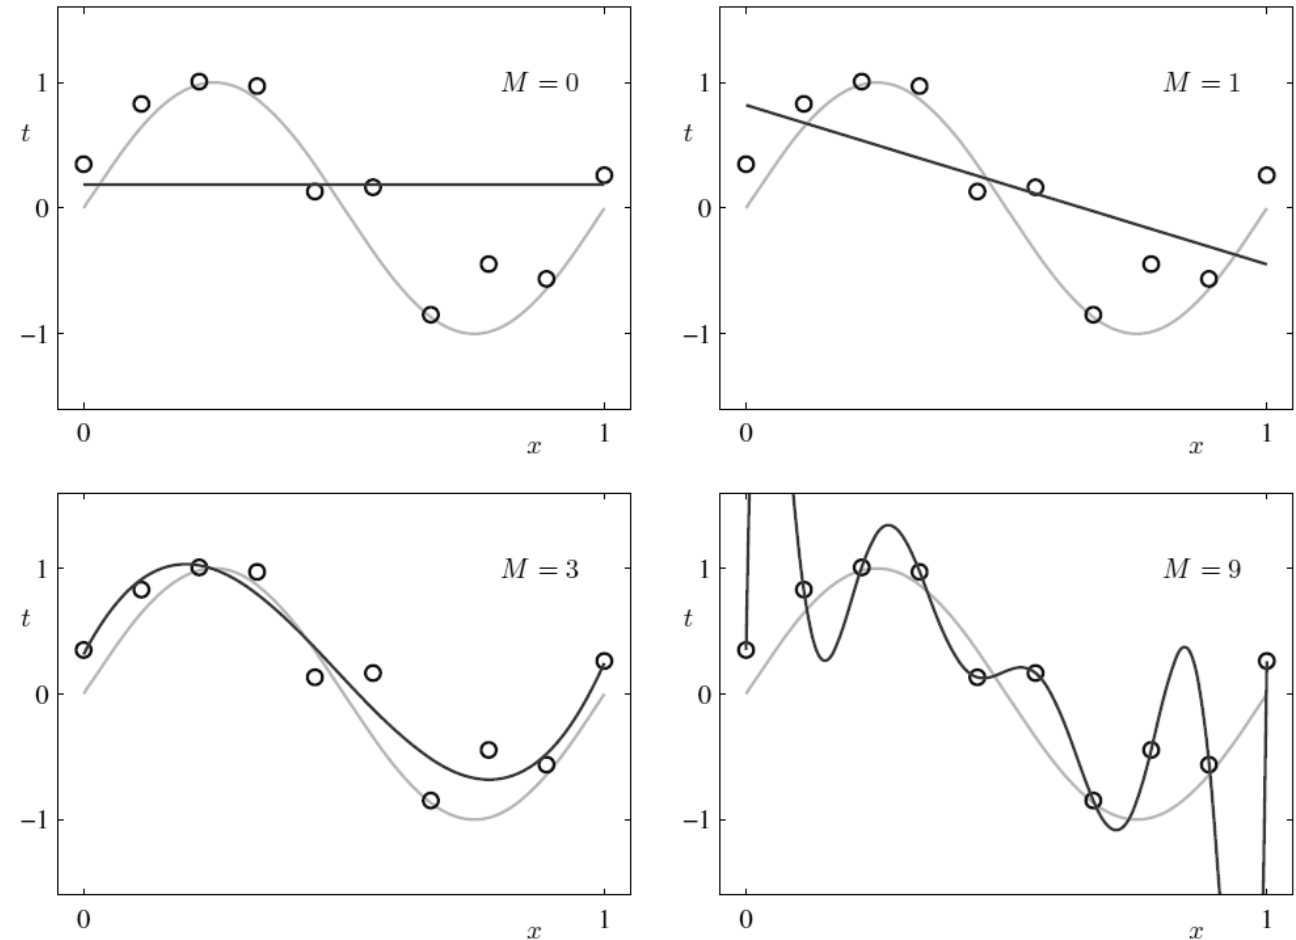
Regularisation

Model complexity, capacity, expressibility, generalisability

Example:

$$y = f(x; \theta) = w_3x^3 + w_2x^2 + w_1x + w_0$$

Overfitting and underfitting



Regularisation

Linear polynomial: $y = f(x; \mathbf{w}) = \sum_{m=0}^M w_m x^m$

Mean-square-error (MSE) as loss: $\ell_{\theta} = \frac{1}{N} \sum_{n=1}^N (y_n - t_n)^2$

L²-Norm (weight-decay): $\|\mathbf{w}\|^2 = \sum_{m=0}^M w_m^2 = w_0^2 + w_1^2 + w_2^2 + \dots$




Hyperparameter: λ

Regularised loss: $\tilde{\ell}_{\theta} = \ell_{\theta} + \frac{\lambda}{2} \|\mathbf{w}\|^2$

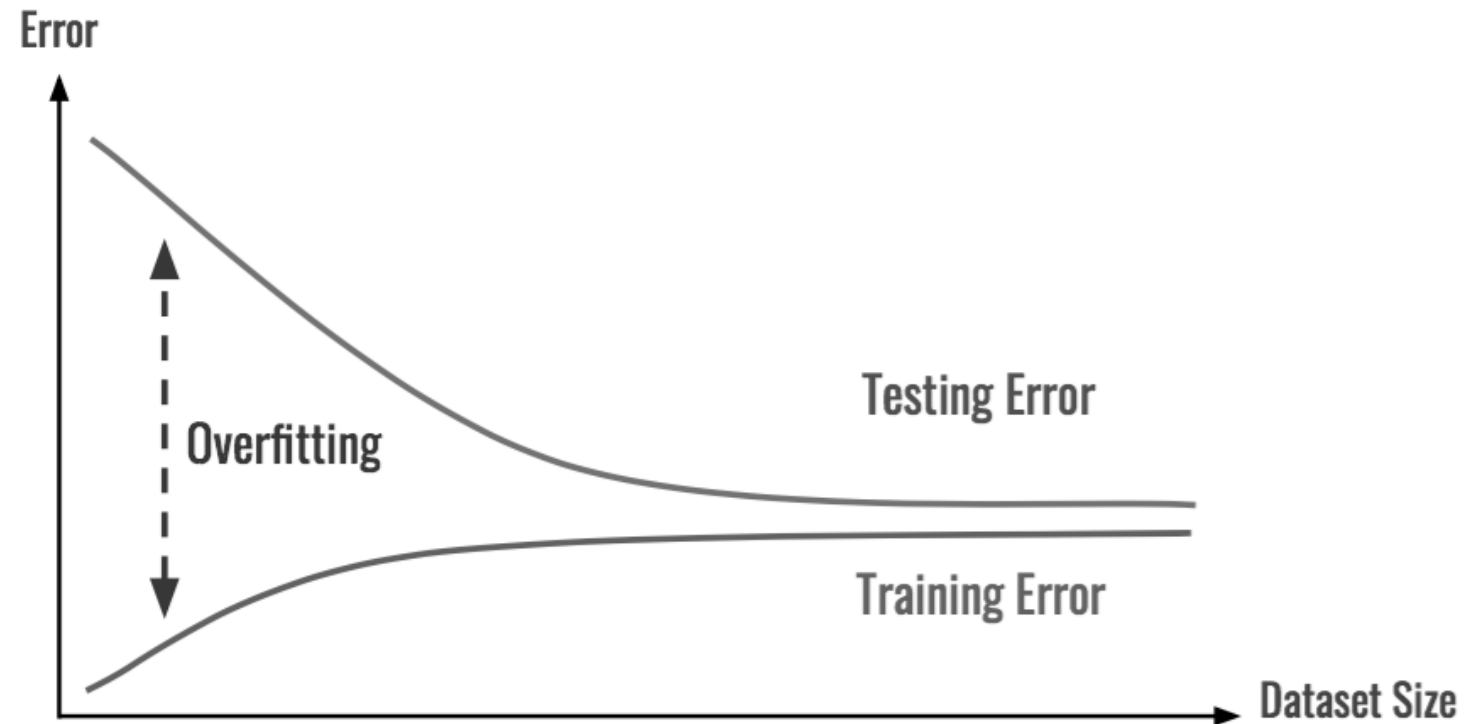
Solution

Least-square solution

Gradient-descent



Generalisability vs. dataset size



- Purpose of regularisation
- Approaches to regularising deep neural networks

- Training strategies

Early stopping, curriculum learning, data resampling.

- Data augmentation

Random data transformation, affinity and diversity

- Invariance and normalisation

Spatial transformer networks, batch normalisation, reparameterization

- Parameter constraints

Parameter norms, sparse representation, parameter sharing, multi-task, semi-supervised

- Unsupervised learning

Autoencoder, generative adversarial network

- Randomness


Noise, dropout

- Model combining

Model averaging, ensemble, bagging/bootstrap aggregating, boosting

Regularisation | Training Strategies


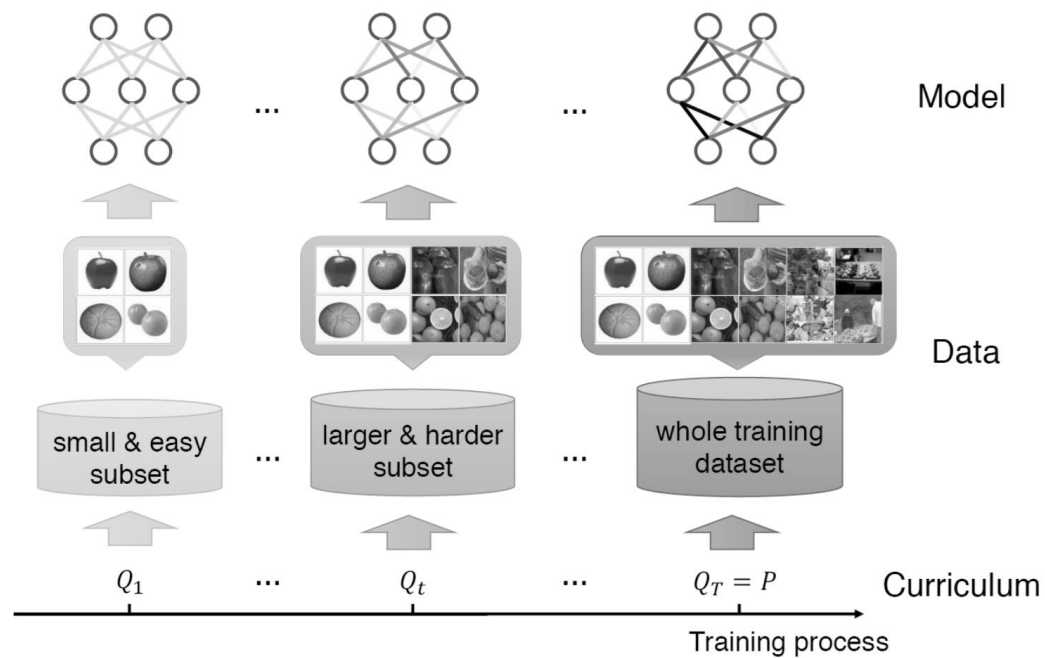
- Early stopping



– Curriculum learning

1. Sorting example difficulty (task complexity*)
2. Pacing the curriculum learning

Example algorithms: data-resampling



– Data resampling

- Stochastic (minibatch) gradient descent, sampling without replacement
- Data resampling, e.g. up/down-sampling for imbalanced classes, similar to weighted loss $\ell_{\mathbf{w}}(\mathbf{y}_n, \mathbf{t}_n; \omega_n) = \frac{1}{N} \sum_{n=1}^N \omega_n d(\mathbf{y}_n, \mathbf{t}_n)$

– Empirical risk minimisation

- Data (generating) distribution: $\mathbb{E}_{(\mathbf{x}, y) \sim p_{data}} [\ell(f(\mathbf{x}, \boldsymbol{\theta}), y)]$
- Training data distribution: $\mathbb{E}_{(\mathbf{x}, y) \sim \hat{p}_{data}} [\ell(f(\mathbf{x}, \boldsymbol{\theta}), y)]$
- Empirical risk, e.g. minibatch: $\mathbb{E}_{(\mathbf{x}, y) \sim \hat{p}_{data}(\mathbf{x}, y)} [\ell(f(\mathbf{x}, \boldsymbol{\theta}), y)] = \frac{1}{M} \sum_{m=1}^M \ell(f(\mathbf{x}^m, \boldsymbol{\theta}), y^m)$

$$\text{i.e. } \hat{p}_{data}(\mathbf{x}, y) = \frac{1}{M} \sum_{m=1}^M \delta(\mathbf{x} = \mathbf{x}^m, y = y^m)$$

- Approximate data distribution using vicinity distribution: $\hat{p}_v(\mathbf{x}, y) = \frac{1}{M} \sum_{m=1}^M v(\tilde{\mathbf{x}}, \tilde{y} | \mathbf{x}^m, y^m)$

$$\text{e.g. } v(\tilde{\mathbf{x}}, \tilde{y} | \mathbf{x}^m, y^m) = \mathcal{N}(\tilde{\mathbf{x}} - \mathbf{x}^m) \delta(y = y^m)$$

– Test data distribution*


– “mixup”

$$\hat{p}_{data}(\mathbf{x}, y) = \frac{1}{M} \sum_{m=1}^M \delta(\mathbf{x} = \mathbf{x}^m, y = y^m)$$

$$\hat{p}_v(\mathbf{x}, y) = \frac{1}{M} \sum_{m=1}^M v(\tilde{\mathbf{x}}, \tilde{y} | \mathbf{x}^m, y^m) = \frac{1}{M} \sum_{m=1}^M \mathcal{N}(\tilde{\mathbf{x}} - \mathbf{x}^m) \delta(y = y^m) \text{ - adding Gaussian noise to data}$$

$$\hat{p}_\mu(\mathbf{x}, y) = \frac{1}{M} \sum_{m=1}^M \mu(\tilde{\mathbf{x}}, \tilde{y} | \mathbf{x}^m, y^m),$$


where $\mu(\tilde{\mathbf{x}}, \tilde{y} | \mathbf{x}^m, y^m) = \frac{1}{N} \sum_{n=1}^N \mathbb{E}[\delta(\tilde{\mathbf{x}} = \lambda \mathbf{x}^m + (1 - \lambda) \mathbf{x}^n, \tilde{y} = \lambda y^m + (1 - \lambda) y^n)]$ - minimising empirical vicinal risk, i.e. mixup



$$\tilde{\mathbf{x}} = \lambda \mathbf{x}^m + (1 - \lambda) \mathbf{x}^n$$

$$\tilde{y} = \lambda y^m + (1 - \lambda) y^n$$

Image



Label

[1.0, 0.0]
cat dog


[0.0, 1.0]
cat dog

[0.7, 0.3]
cat dog

– Data augmentation

Regularisation | Data Augmentation

- Random data transformation
 - Colour/intensity/contrast space
 - Imaging-specific parameters
 - e.g. photometric transformation (luminance, illuminance, flux, intensity...),
 - bias field for MR,
 - perspective transformation
 - Spatial transformation
 - Geometric: flipping, cropping, rotation
 - Affine
 - Nonlinear deformation
- (Unsupervised) generative models*



– Affinity and diversity


- Affinity quantifies how much an augmentation shifts the training data distribution from that learned by a model.
e.g. difference in validation accuracies due to augmentation

$$\mathcal{T}[a; m; D_{val}] = \mathcal{A}(m, D'_{val}) - \mathcal{A}(m, D_{val}) .$$


- Diversity quantifies the complexity of the augmented data with respect to the model and learning procedure.
e.g. expected training loss due to augmentation

$$\mathcal{D}[a; m; D_{train}] := \mathbb{E}_{D'_{train}} [L_{train}] .$$

– Other definitions*



(a) Affinity vs Diversity

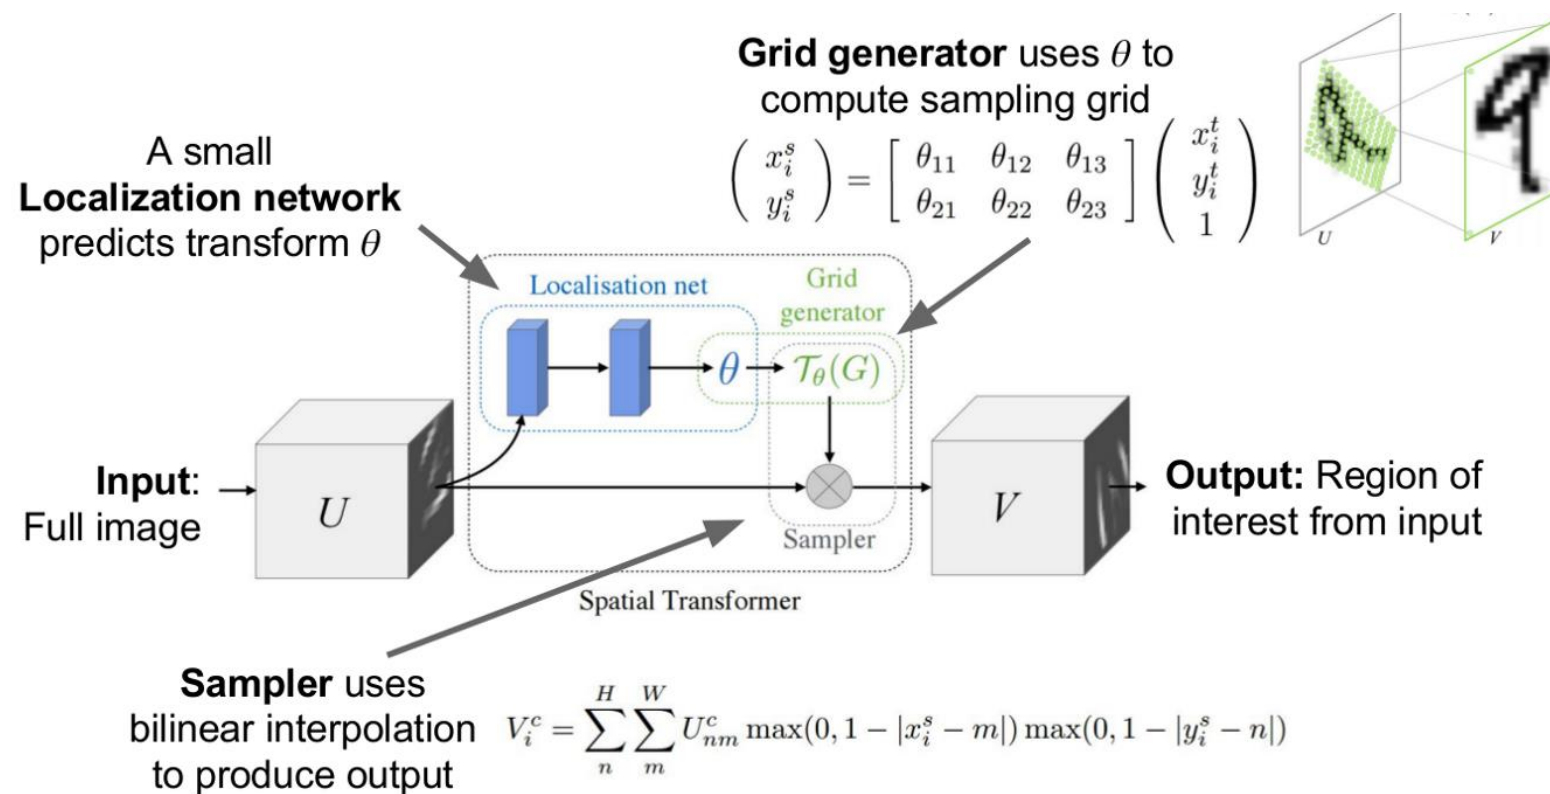


(b) Model's View of Data

Regularisation | Invariance and Normalisation

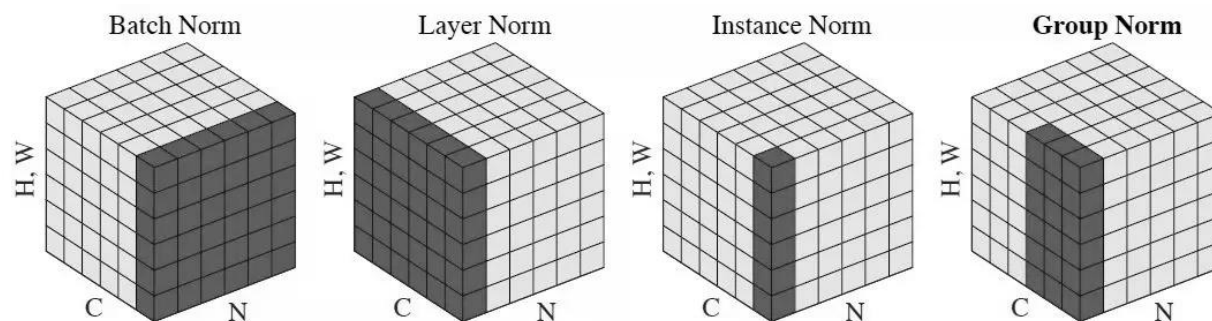
– Spatial transformer networks

“Random transformation in training a model = encouraging the model invariant to the transformation”



– Batch normalisation

- Normalise features to a standard Normal distribution within mini-batch
- Learnable parameters for linearly transforming maintaining expressibility
- Makes bias redundant in previous network layers
- Use population statistics during inference
- Benefits
 - Introduce both random additive and multiplicative noise during training
 - Reduce inter-layer dependency



- Batch/layer/instance/group normalization
- BN for CNN and LN for RNN?
- A form of re-parameterisation of layer activations

$$\hat{\mathbf{x}} = \frac{\mathbf{x} - \mu}{\sigma}$$

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

- Weight normalisation

$$\mathbf{w} = \frac{g}{\|\mathbf{v}\|} \mathbf{v}$$

- Input normalisation
- Label normalisation – data scaling
- Permutation invariance
- Translation invariance
- ...

Regularisation | Parameter Constraints

- Parameter norms

- Weight decay by penalising L2- and L1-norms

$$\ln p(\mathbf{w}|\mathbf{t}) = -\frac{\beta}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 - \frac{\alpha}{2} \mathbf{w}^T \mathbf{w} + \text{const.}$$

- Difference between L2- and L1-norms

$$\Omega(\boldsymbol{\theta}) = \|\mathbf{w}\|_1 = \sum_i |w_i| \qquad \nabla_{\mathbf{w}} \tilde{J}(\mathbf{w}; \mathbf{X}, \mathbf{y}) = \alpha \text{sign}(\mathbf{w}) + \nabla_{\mathbf{w}} J(\mathbf{X}, \mathbf{y}; \mathbf{w})$$

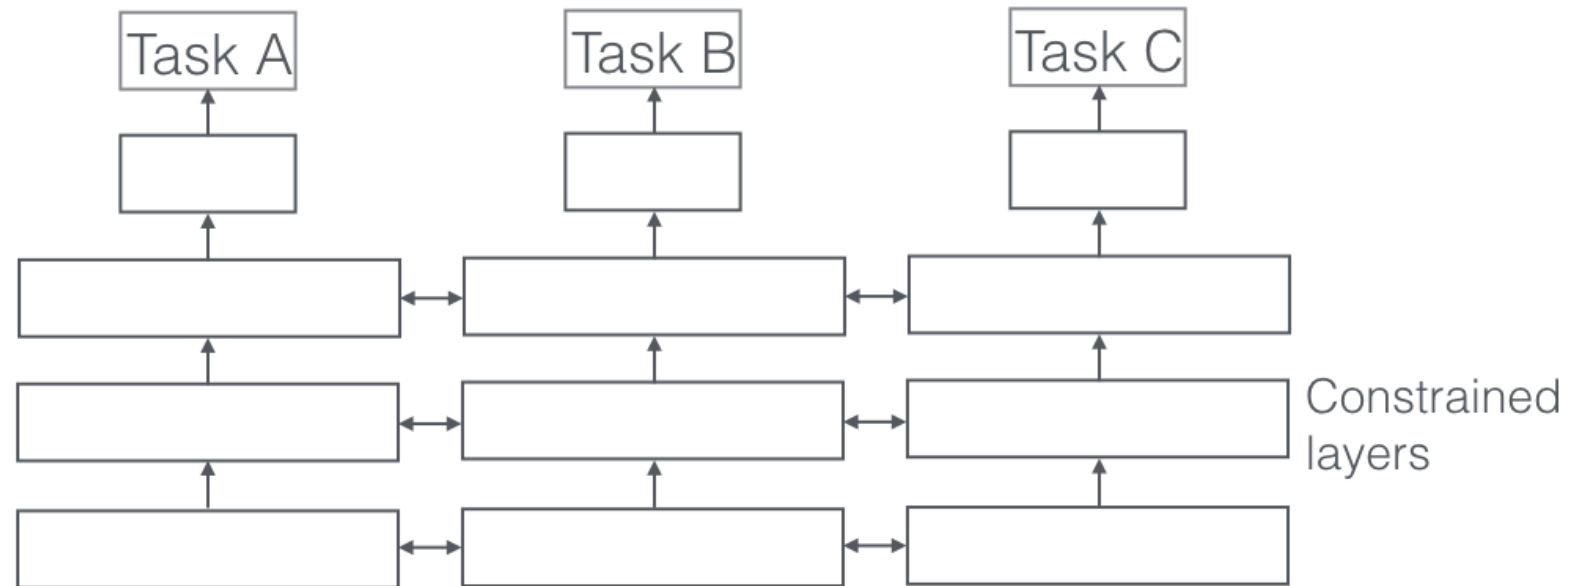
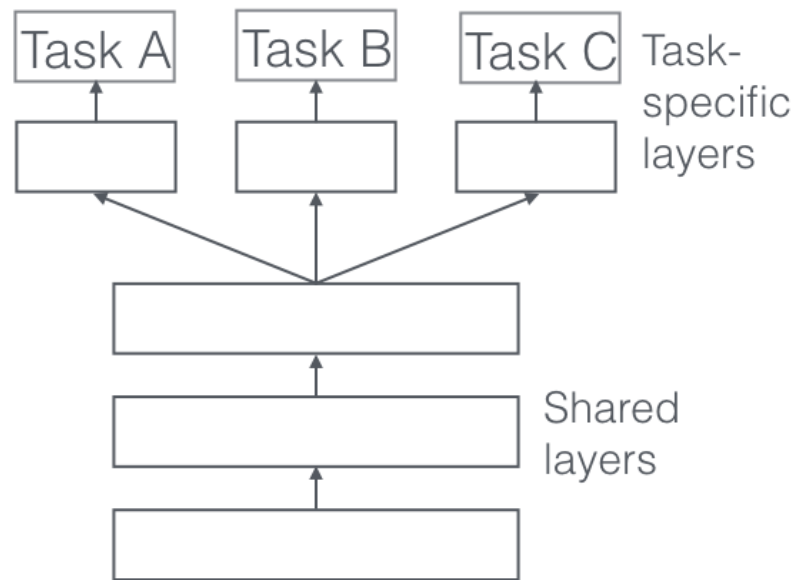
- L2 norm gradient scales with w, therefore less likely to become zero, i.e. sparsity.

- Parameter sharing
- CNN
- RNN
- Tiling

$$\Omega(\boldsymbol{w}^{(A)}, \boldsymbol{w}^{(B)}) = \|\boldsymbol{w}^{(A)} - \boldsymbol{w}^{(B)}\|_2^2$$

- Multi-task learning*
- Semi-supervised learning*.


- Multi-task learning



- Multi-task learning

- Generalised formulation


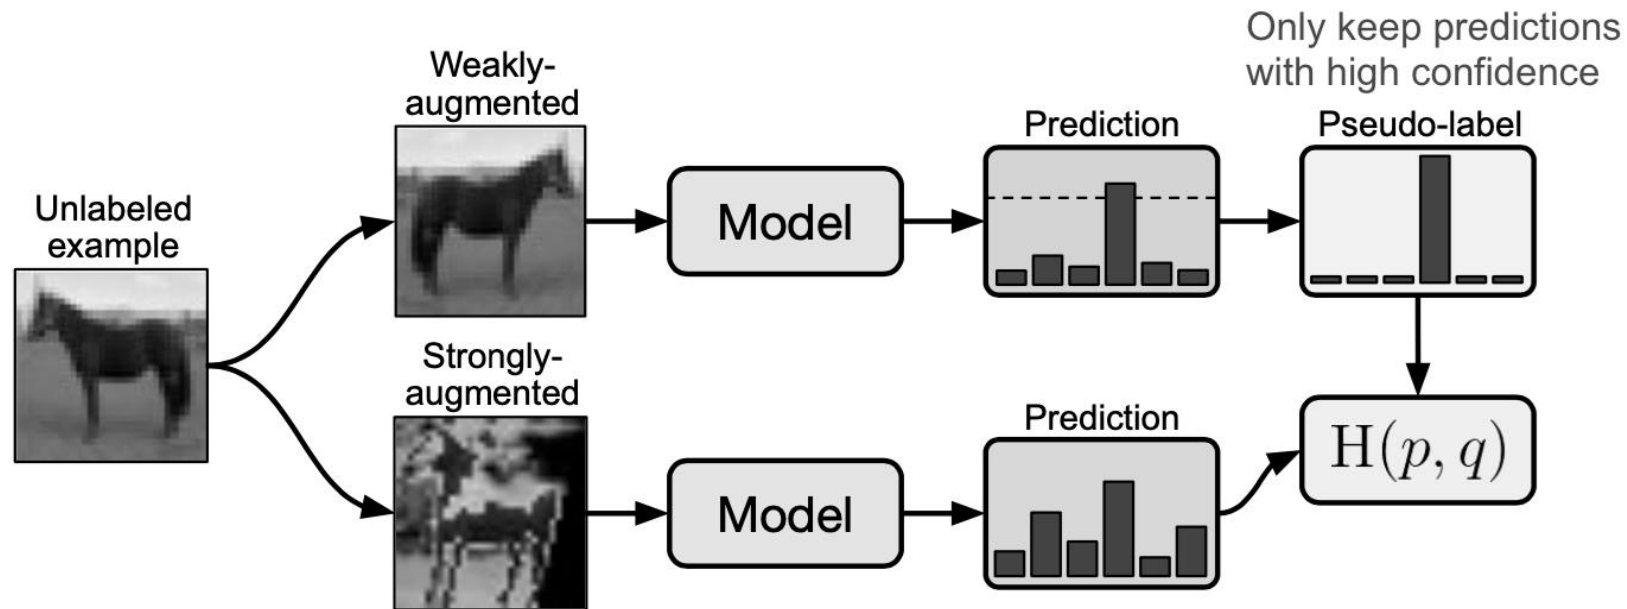
- Z_i : One-hot task index / indicator / descriptor
- $f(Y|X) \rightarrow f(Y|X, Z_i)$



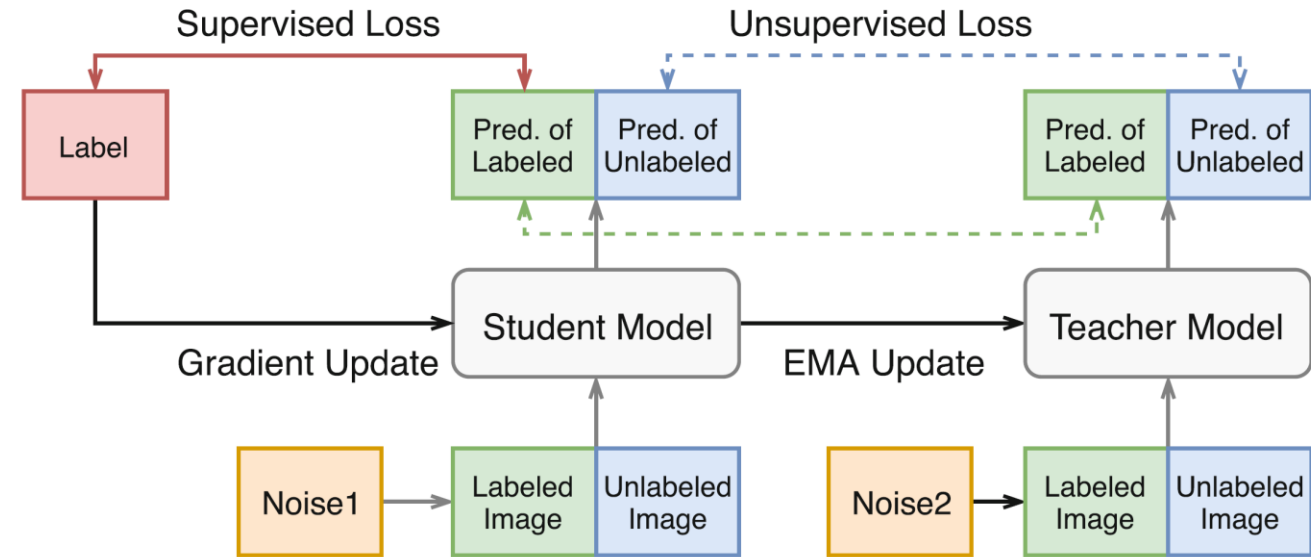
- Implementation

- Shared and task specific parameters (negative transfer)
- Concatenation/summation
- Multiplication conditioning (gating/multi-head)

- Semi-supervised learning
- Supervised learning: labelled data = data + labels, $f(Y|X)$
- Semi-supervised learning: labelled data $f(X, Y)$ + unlabelled data $f(X)$
- Pseudo labels and entropy minimization, consistency



- Semi-supervised learning
 - Consistency between data augmented predictions
 - MixedMatch
 - Student-Teacher




Regularisation | Unsupervised Learning

Autoencoder

Autoencoder architecture: encoder - (low-dimension) "code" - decoder

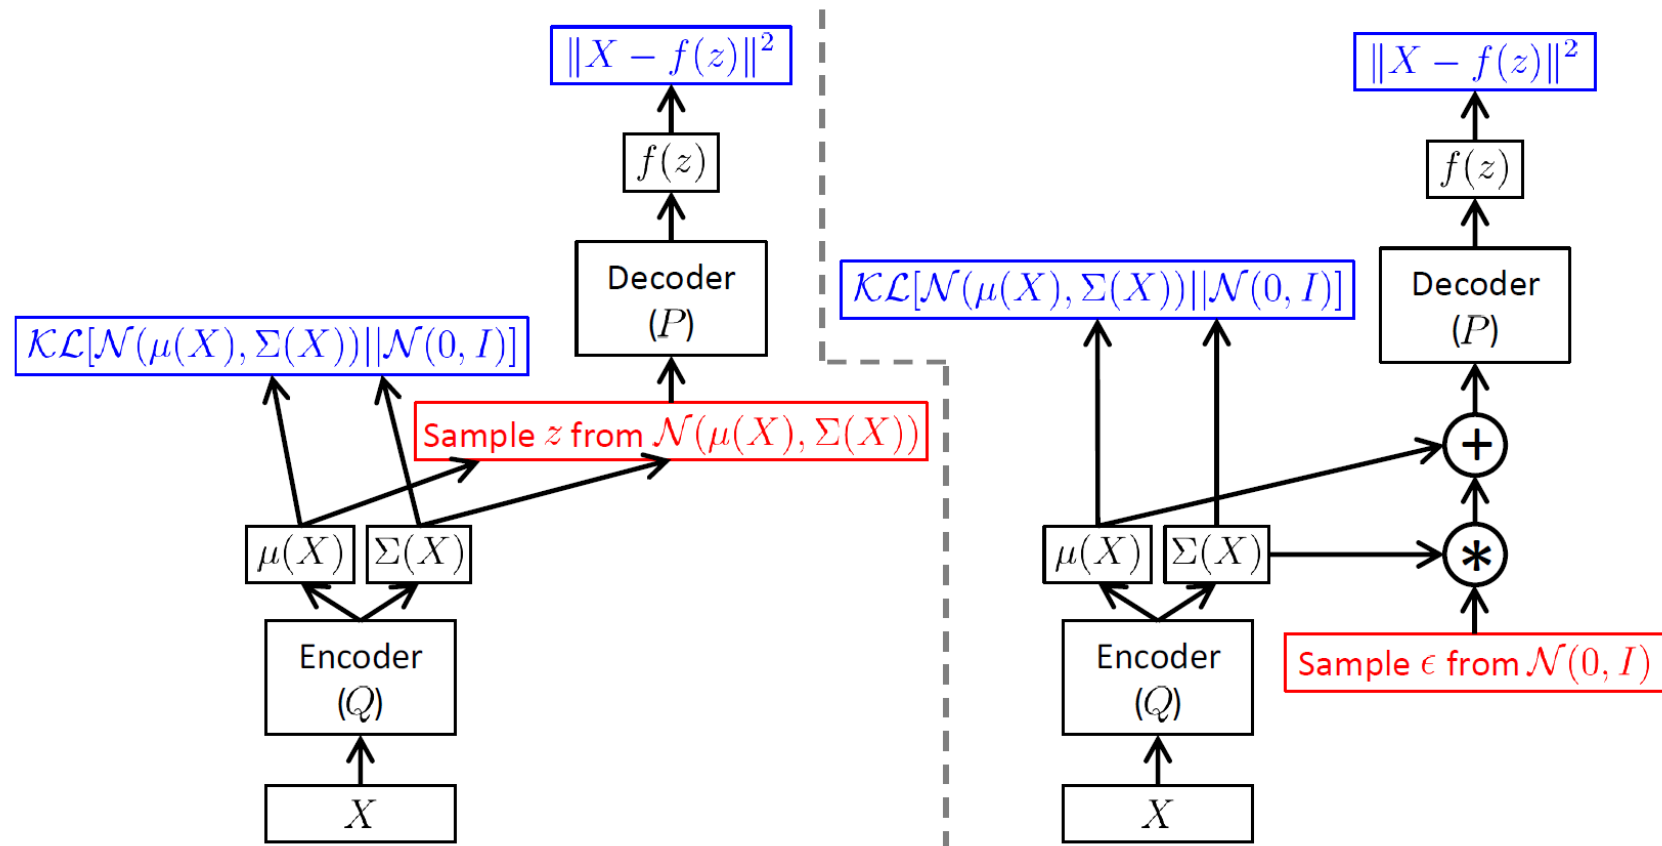
Training loss: self-reconstructing difference, e.g. CE, MSE

Latent "code space"

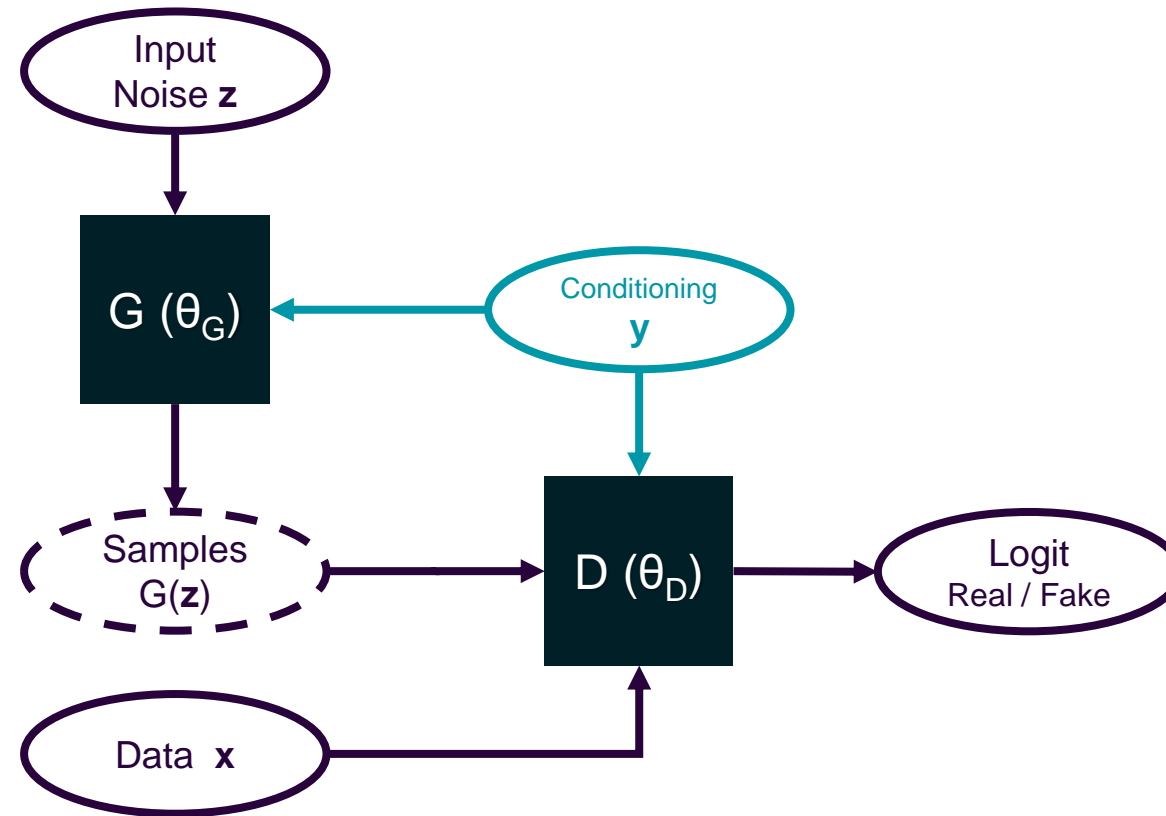


Variational autoencoder

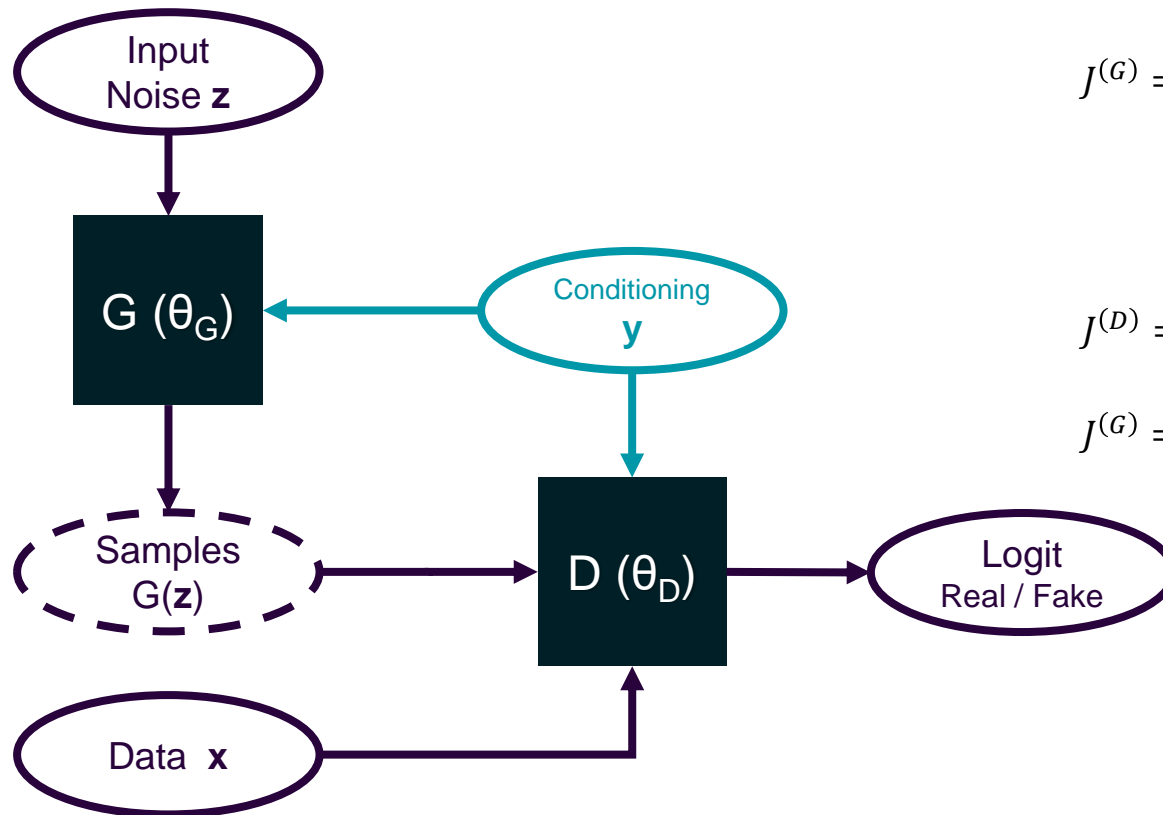
Reparameterisation of latent space



Generative adversarial networks



Generative adversarial networks



$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{\mathbf{x} \sim P_{data}} \log D(\mathbf{x}) - \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim N} \log (1 - D(G(\mathbf{z})))$$

$$J^{(G)} = -\frac{1}{2} \mathbb{E}_{\mathbf{z} \sim N} \log D(G(\mathbf{z}))$$

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim P_{data}} \log D(\mathbf{x}, \mathbf{y}) - \frac{1}{2} \mathbb{E}_{\mathbf{z} \sim N, \mathbf{y} \sim P_{cond}} \log (1 - D(G(\mathbf{z}, \mathbf{y}), \mathbf{y}))$$

$$J^{(G)} = -\frac{1}{2} \mathbb{E}_{\mathbf{z} \sim N, \mathbf{y} \sim P_{cond}} \log D(G(\mathbf{z}, \mathbf{y}), \mathbf{y})$$

Regularisation | Randomness

Noise

- Gaussian
- Uniform with predefined range
- Bernoulli distribution, i.e. dropout


- Inputs, e.g. augmentation, invariant to noise
- Weights, e.g. RNN, at initialisation
- Output targets, e.g. label smoothing
- Edges / hidden units, e.g. dropout*

- Sampling from uncertainty
- Equivalent to penalty on weight norms


- ...

Dropout

- Drop as many as 80% of nodes
- Use all nodes at inference
- Dropout as model ensemble
- Dropout as random noise
- Dropout as Bayesian model sampling
- Uncertainty estimation using dropout
- Posterior estimation at inference
- With convolutional layers?



(a) Standard Neural Net

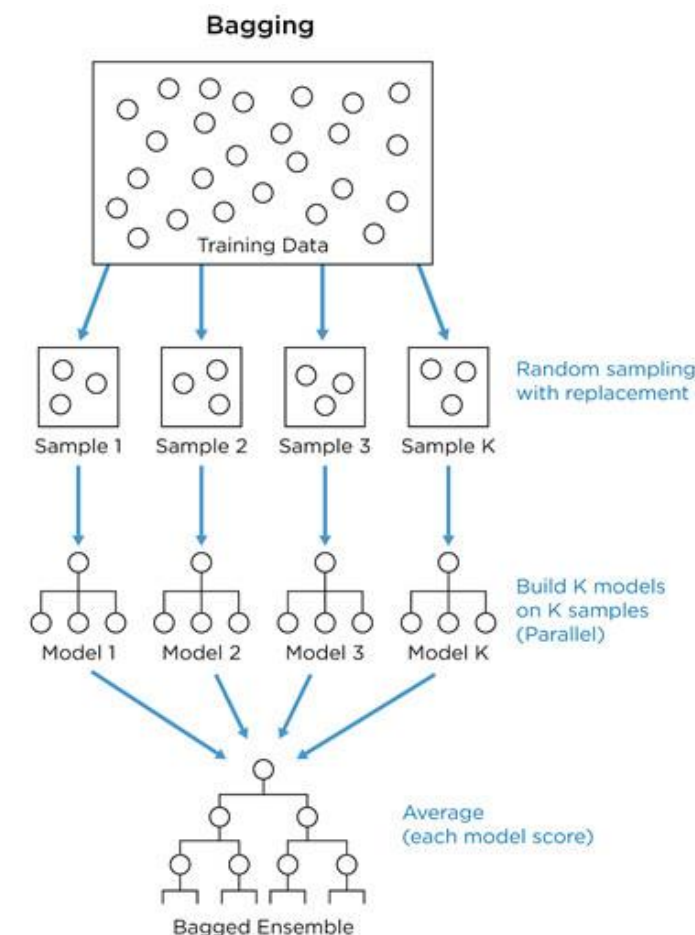


(b) After applying dropout.

Regularisation | Model Combining

Bagging = bootstrap aggregating \in model averaging = ensemble?

- Independent errors are multiplicative
- Fully correlated errors does not accumulate
- Optimising multiple low-bias models “cancels” variance, i.e. more samples to estimate
- Bagging: the same kind of model, trained multiple times with bootstrapped datasets
 - Re-trained on the same dataset
 - Challenge winners!
- Committee: multiple models
 - e.g. dropout
- Boosting: multiple models in sequence (convert weak learners to a strong learning)
 - e.g. Tree-based models*



- Training strategies

Early stopping, curriculum learning, data resampling.

- Data augmentation

Random data transformation, affinity and diversity

- Invariance and normalisation

Spatial transformer networks, batch normalisation, reparameterization

- Parameter constraints

Parameter norms, sparse representation, parameter sharing, multi-task, semi-supervised

- Unsupervised learning


Autoencoder, generative adversarial network


- Randomness

Noise, dropout

- Model combining

Model averaging, ensemble, bagging/bootstrap aggregating, boosting





Add two different regularization methods to the “image classification” tutorial and compare the results.